

Normalization Without Convergence: Large Language Models, Linguistic Entropy, and the Myth of Cognitive Homogenization

Flyxion

Abstract

Recent commentary in cognitive science and computational linguistics has proposed that large language models (LLMs) are driving a broad homogenization of human language, perspective, and reasoning. According to this view, the widespread use of generative systems encourages convergence toward a statistical “median” embedded in the models’ training distributions, potentially narrowing the diversity of linguistic expression and epistemic viewpoints across communicating populations. While this hypothesis raises important questions about the long-run cultural consequences of AI-mediated communication, many existing formulations rest on an implicit and under-examined assumption: that linguistic output provides a transparent window into the underlying generative structure of human cognition.

This paper argues that such claims systematically conflate two distinct phenomena: *surface-level normalization in the communication channel* and *deep convergence in the generative processes of human cognition*. Drawing on an information-theoretic framework, we demonstrate that a system capable of mapping many idiosyncratic inputs to a standardized representation necessarily reduces the entropy of observable outputs while leaving the entropy of the underlying source distribution unchanged or unaffected. Formally, this follows from the data-processing inequality: post-processing cannot increase mutual information. In other words, uniformity at the level of representational output does not entail—and cannot establish—homogenization at the level of thought.

From this perspective, large language models function less as normative enforcement mechanisms than as *translation layers* between heterogeneous private linguistic systems and shared public registers. By reducing the communicative cost of nonstandard spelling, dialect, or multilingual composition, such systems may in fact expand the effective expressive space available to individual authors. We formalize this argument through a source-channel

model, develop the notion of *linguistic attractor basins* in a probabilistic landscape, and introduce a distinction between the entropy of generative sources and the entropy of exploratory discovery processes.

We further argue that the more significant ecological risk of LLM-mediated communication lies not in stylistic convergence but in the potential contraction of *exploratory variance* in knowledge discovery: the diversity of trajectories through which individuals encounter novel ideas, conflicting interpretations, and unexpected conceptual recombinations. If model-mediated search increasingly directs users toward the same conceptual summaries and canonical explanations, the space of intellectual trajectories may contract even while expressive diversity at the surface level remains intact. Understanding the cultural consequences of generative AI therefore requires attention not only to the entropy of linguistic form but also to the information ecology of exploration, discovery, and conceptual variation.

Contents

1	Introduction	7
2	The Homogenization Thesis	9
2.1	Linguistic Convergence	9
2.2	Perspectival Convergence	10
2.3	Reasoning Convergence	10
2.4	The Shared Inferential Structure	10
3	Source Entropy and Channel Normalization	11
3.1	The Communication Model	11
3.2	Entropy of the Generative Source	11
3.3	Entropy of the Normalized Representation	12
3.4	The Data Processing Inequality	13
3.5	Formal Implications for the Homogenization Thesis	13
4	Private Generative Systems and Public Registers	14
4.1	Historical Normalization Mechanisms	14
4.2	LLMs as Automated Normalization Layers	15
4.3	The Expansion of Expressive Space	15
5	Instrumental and Epistemic Interaction	16
5.1	Instrumental Use	16
5.2	Epistemic Use	16
5.3	The Importance of the Distinction	17
6	Linguistic Attractor Basins	17
7	The Conceptual Landscape	18
7.1	Attractor Dynamics Under Repeated Interaction	19
8	Conceptual Trajectories and Intellectual Identity	20
8.1	Conceptual Space and the Topic Simplex	20
8.2	Conceptual Trajectories	21
8.3	Trajectory Distance and Author Divergence	21
8.4	Topic Constellations	22
8.5	Scaling of Conceptual Signatures	22
9	Topic Monotonicity and Conceptual Toolchains	23
9.1	Conceptual Toolchains	23
9.2	Topic Monotonicity	24
9.3	Normalization and Toolchain Preservation	24

9.4	Intellectual Identity as a Structured Conceptual Toolkit	25
10	Multimodal Signatures and Input-Layer Diversity	25
10.1	Behavioral Signals in Communication	25
10.2	Generative Model of Expression	26
10.3	Multimodal Authorship Signatures	27
10.4	Scaling Diversity Through Input Channels	27
11	Dialect Liberation Under Automated Normalization	28
11.1	Communication Friction and Standardization Pressure	28
11.2	Automated Normalization as a Cost-Reducing Layer	29
11.3	Expansion of Private Linguistic Systems	29
11.4	Decoupling of Generative and Communicative Registers	30
11.5	Implications for Linguistic Diversity	30
12	Constraint Geometry of Conceptual Spaces	31
12.1	Normalization as Projection onto the Linguistic Manifold	31
12.2	Strata and Cross-Stratum Transition	32
13	Expressive Degeneracy in Linguistic Systems	32
13.1	Many-to-One Structure of Expressive Mapping	33
13.2	Implications for Interpretability of Output Similarity	33
13.3	Degeneracy Grows with Dimensionality	34
14	Combinatorial Expansion in Conceptual Space	34
14.1	Combinatorial Structure of Topic Combinations	35
14.2	Coordination of Representation and Expansion of Content	35
14.3	Innovation from Rare Combinations	35
15	Transverse Innovation in Conceptual Manifolds	36
15.1	Longitudinal versus Transverse Intellectual Movement	36
15.2	Independence of Transverse Movement from Stylistic Form	37
15.3	Preservation of Transverse Directions	37
16	Exploratory Entropy and the Ecology of Knowledge	37
16.1	Exploratory Entropy Defined	38
16.2	Traditional Knowledge Environments	38
16.3	Algorithmic Mediation and the Concentration of Discovery	39
16.4	The Significance of Exploratory Variance	40

17 Rank-Mediated Discovery and Its Entropy	40
17.1 Ranking Functions and Exposure Distributions	40
17.2 Compounding Effects of Ranking and Summarization	41
17.3 Design as a Variable	41
18 Normalization Layers in Historical Communication Systems	42
18.1 Case Studies in Communicative Standardization	42
18.2 The Structural Pattern	43
19 Reexamining Recent Empirical Claims	43
19.1 Representation Versus Generation	43
19.2 Instrumental Versus Epistemic Interaction	44
19.3 Short-Term Effects and Long-Term Dynamics	44
19.4 Stylistic Similarity and Conceptual Diversity	44
20 Toward a Precise Research Agenda	45
20.1 Longitudinal Studies of Generative Behavior	45
20.2 Transfer Studies in Reasoning	45
20.3 Measurement of Exploratory Entropy	45
21 Cognitive Prostheses and Human Agency	46
21.1 The Prosthetic Tradition	46
21.2 LLMs as Prosthetic Systems	46
21.3 Human Agency in the Prosthetic System	47
22 Implications for Linguistic Diversity	47
22.1 Normalization and Diversity Are Compatible	47
22.2 Normalization May Expand Effective Expressive Space	48
22.3 The Ecological Risk Is Real but Differently Located	48
22.4 Design Choices Determine Cultural Consequences	48
23 Externalization of Linguistic Effort and the Redistribution of Cognitive Load	48
23.1 Reallocation of Compositional Resources	49
23.2 The Separation of Generation and Representation	49
23.3 Historical Precedents	50
24 Temporal Signatures in Intellectual Production	50
24.1 Temporal Structure of Conceptual Trajectories	50
24.2 Temporal Signatures as Author Fingerprints	51
24.3 Temporal Records Under Reduced Friction	51

25 A Multilayer Model of Expressive Diversity	52
25.1 The Five Layers	52
25.2 Layer Independence and the Invalidity of Homogenization Inferences	53
25.3 Implications for Research and Policy	54
26 Conclusion: Diversity, Compression, and Cultural Dynamics	54
A Formal Derivation of Entropy Inequalities Under Deterministic Normalization	56
A.1 Setup	56
A.2 Entropy of the Source	56
A.3 Entropy of the Normalized Output	56
A.4 The Core Inequality	56
A.5 Quantifying Information Loss	57
A.6 The Data Processing Inequality	57
B Exploratory Entropy: Formal Treatment	57
B.1 Setup	57
B.2 Maximum Entropy and the Uniform Distribution	58
B.3 Entropy Under Ranking	58
B.4 Entropy Under Summarization	58
C A Unified Proposition on Normalization and Exploration	59

1 Introduction

Technologies that mediate human communication have repeatedly and profoundly reshaped the structure of linguistic expression and intellectual exchange. Writing systems, printing presses, editorial institutions, telegraph networks, search engines, and digital archives have each altered the mechanisms through which language is produced, transmitted, and interpreted. These transformations have recurrently generated concerns about the potential homogenization of thought: that the machinery of communication will ultimately flatten the diversity of the minds it connects. Yet the historical record consistently resists this conclusion. The diffusion of moveable type standardized orthographic conventions across European languages without destroying the intellectual heterogeneity of the Reformation. Editorial institutions imposed consistent rhetorical styles on scientific writing without constraining the diversity of scientific ideas. The internet concentrated written discourse around a small number of dominant platforms without eliminating the extraordinary proliferation of viewpoints that those platforms carry. In each case, normalization in the communication channel coexisted with the persistence or even expansion of cognitive diversity at the level of the source.

The emergence of large language models (LLMs) represents the latest development in this long sequence of communicative technologies. Generative AI systems are now widely used to assist with writing, translation, summarization, coding, and information retrieval. Because these systems produce text by sampling from probability distributions learned from large corpora, their outputs often reflect statistically common linguistic forms: frequently recurring syntactic constructions, standard rhetorical organization, and stylistically moderate registers. This characteristic has prompted a growing body of scholarly commentary suggesting that the widespread adoption of LLMs may gradually homogenize human language, perspective, and reasoning—that individual discourse will be pulled, through repeated interaction with shared systems, toward a common statistical center.

The homogenization hypothesis raises important questions regarding the long-term cultural consequences of AI-mediated communication. If a large and growing proportion of human discourse becomes filtered through a small set of generative systems, the observable distribution of linguistic expression may indeed shift toward a common mode. Such a shift could plausibly influence not only the stylistic surface of language but the structure of argumentation and the framing of ideas. These are legitimate concerns that deserve serious analytical treatment.

However, evaluating the homogenization hypothesis requires careful attention to a distinction that much of the existing literature has failed to draw clearly: the distinction between *observable representations* and the *generative processes* that produce them. Communication systems routinely include intermediate transfor-

mation layers that reshape linguistic signals during transmission without altering the cognitive diversity of the communicating agents. Spelling conventions normalize orthographic variation. Editorial practices standardize rhetorical structure. Translation systems project expressions across language boundaries. Each of these mechanisms compresses surface-level variation without penetrating the generative layer of cognition.

Large language models can be interpreted as a technologically novel form of such normalization. When used as editing or translation tools, they map diverse linguistic inputs into standardized outputs. The resulting reduction in observable variability may therefore reflect the operation of the communication channel rather than any transformation of the generative source. If this interpretation is correct, the inference from observed output uniformity to cognitive convergence is formally invalid.

This paper develops a theoretical framework for analyzing this distinction with precision, drawing on concepts from information theory, probabilistic modeling, and the ecology of knowledge systems. The central claim is that the homogenization thesis systematically conflates two distinct kinds of entropy reduction: compression within the communication channel, which is a property of the transformation, and convergence within the generative processes of human cognition, which is a claim about the source distribution.

The analysis is organized as follows. Section 2 introduces the homogenization hypothesis in its principal variants and identifies the shared inferential assumptions that underlie them. Section 3 develops an information-theoretic framework distinguishing source entropy from channel entropy and establishes the central formal results. Section 4 elaborates the distinction between private generative systems and public communicative registers, drawing on historical precedent. Section 5 analyzes the difference between instrumental and epistemic uses of generative models and their differential implications for cognitive influence. Sections 6–7 develop the geometric notion of linguistic attractor basins and the conceptual landscape in which they reside. Sections 8–15 extend this geometric framework to analyze conceptual trajectories, topic monotonicity, multimodal signatures, dialect liberation, constraint geometry, expressive degeneracy, combinatorial expansion, and transverse innovation across conceptual strata. Section 16 introduces the concept of exploratory entropy and analyzes how algorithmic mediation reshapes discovery distributions. Sections 19 and 20 engage with the specific claims of recent empirical literature and propose a more precise research program. Sections 21–25 situate LLMs within a broader history of cognitive prosthetics, analyze the externalization of linguistic labor, examine temporal signatures in intellectual production, and synthesize the analysis into a unified five-layer model of expressive diversity. Section 26 draws the threads together and states the paper’s principal conclusions.

Appendices provide more detailed derivations of the central entropy results and a formal treatment of rank-mediated discovery distributions.

2 The Homogenization Thesis

Recent discussions in cognitive science and computational linguistics have articulated a growing concern that large language models may induce a progressive homogenization of human language, reasoning, and perspective. The core intuition motivating this family of claims is straightforward. Because LLMs are trained on large corpora that aggregate statistical regularities across vast quantities of human communication, their outputs tend to reflect the most probable linguistic forms and conceptual framings present in those datasets. When such systems are widely adopted as writing assistants, editors, and conversational partners, it is argued that human discourse may gradually converge toward the same statistical center that those systems implicitly represent [Sourati et al., 2026].

In this view, the widespread use of generative language models acts as a powerful attractor within the space of linguistic and epistemic practices. Individual expressions that initially vary across dialects, registers, educational backgrounds, and cultural contexts are progressively “pulled” toward a shared normative distribution. Over time, this process is thought to reduce variance across several dimensions of human communication simultaneously: lexical choice, rhetorical structure, argumentative form, and epistemic framing.

The homogenization thesis appears in three structurally distinct but related variants that are worth distinguishing precisely.

2.1 Linguistic Convergence

The *linguistic convergence* hypothesis proposes that texts produced with the assistance of LLMs increasingly resemble one another in vocabulary, syntax, and discourse structure. Because the models generate outputs according to learned probability distributions over sequences of tokens, repeated use of the same systems may produce recognizable stylistic regularities across otherwise independent authors. The distributional signature of the model bleeds into the distributional signature of the assisted writer.

This form of the hypothesis is the most empirically tractable, since it concerns surface properties of text that can in principle be measured directly. Its weakness, as argued below, is that surface similarity among assisted texts is entirely compatible with high diversity in the underlying cognitive processes that generated the pre-normalized inputs.

2.2 Perspectival Convergence

The *perspectival convergence* hypothesis suggests that LLMs encode the dominant assumptions and interpretive frameworks of their training data and transmit those frameworks to users through the authority that attaches to model-generated outputs. Since large portions of the digital corpus from which these models are trained originate in Western academic, journalistic, and institutional contexts, critics argue that the models may implicitly reinforce the epistemic norms, evaluative standards, and categorical schemes associated with those contexts [Bender et al., 2021]. Users who repeatedly encounter these norms may internalize them as default frameworks for understanding the world.

This variant is more difficult to operationalize empirically, since it requires measuring shifts in epistemic orientation rather than surface linguistic properties. It also makes stronger assumptions about the mechanism of influence, implying that repeated exposure to model outputs reshapes the prior beliefs and conceptual categories of users in measurable ways.

2.3 Reasoning Convergence

The *reasoning convergence* hypothesis proposes that repeated interaction with model-generated explanations may influence the inferential strategies humans employ when solving problems. Step-by-step reasoning formats produced by language models might serve as templates for human inference, gradually reshaping problem-solving architecture toward the heuristics those models are tuned to produce.

This is the strongest form of the hypothesis and also the least supported by available evidence. Demonstrating reasoning convergence would require longitudinal evidence of durable transfer from model-mediated reasoning contexts to independent problem-solving behavior—a much higher empirical bar than demonstrating stylistic similarity in assisted texts.

2.4 The Shared Inferential Structure

Despite their differences, all three variants of the homogenization thesis share a common inferential structure. Each proceeds from an observation about the properties of LLM outputs or LLM-assisted human outputs to a claim about the underlying generative processes of human cognition. They assume, explicitly or implicitly, that observable linguistic behavior provides a reliable proxy for the internal structure of human thought.

This assumption is the central target of the present analysis. The argument that follows does not deny that generative AI systems influence human behavior—

they surely do, as do all tools that mediate communication and cognition. The argument is rather that the specific inference from output uniformity to generative convergence is formally invalid, because a normalization layer in the communication channel can produce the former without the latter.

3 Source Entropy and Channel Normalization

A rigorous analysis of the homogenization hypothesis requires placing the communication process within an information-theoretic framework. In classical information theory, a communication system involves at minimum three components: a source that generates signals, a channel through which those signals are transmitted, and a receiver that interprets the resulting representation [Shannon, 1948]. The statistical properties of what the receiver observes depend on the joint operation of the source and the channel. Crucially, these two contributions need not mirror one another.

3.1 The Communication Model

In the context of LLM-mediated communication, the components of the standard model admit natural interpretations. The *generative source* corresponds to the cognitive and linguistic processes through which a human author formulates and expresses ideas. The *normalization layer* corresponds to the large language model when it is used to edit, rewrite, translate, or otherwise transform the human-generated text. The *communication channel* corresponds to the medium through which the normalized representation is transmitted. The *receiver* corresponds to the audience that interprets the resulting output.

Let X denote a random variable representing expressions drawn from the generative source distribution $p(x)$. Let T denote the deterministic transformation applied by the normalization layer, so that the observable output is $Y = T(X)$. The distribution of Y is given by

$$p(y) = \sum_{\{x:T(x)=y\}} p(x).$$

This is the pushforward of the source distribution under the map T .

3.2 Entropy of the Generative Source

The entropy of the generative source distribution quantifies the diversity of linguistic expressions produced by the human author:

$$H(X) = - \sum_x p(x) \log p(x).$$

High values of $H(X)$ indicate that the author draws from a wide and roughly uniform distribution over possible expressions; low values indicate that the author's choices are heavily concentrated around a small set of forms.

3.3 Entropy of the Normalized Representation

The entropy of the observable representation is

$$H(Y) = - \sum_y p(y) \log p(y).$$

When the transformation T is many-to-one—that is, when multiple distinct inputs $x_1, x_2, \dots$ all map to the same output y —the output distribution $p(y)$ is strictly more concentrated than the source distribution $p(x)$. In this case $H(Y) < H(X)$.

Definition 3.1 (Compression Ratio). The *compression ratio* of a normalization transformation T at a point y in the output space is

$$\kappa(y) = |\{x : T(x) = y\}|.$$

If $\kappa(y) > 1$ for any y , the transformation is non-injective and compresses the source distribution.

Proposition 3.2 (Entropy Inequality). *For any deterministic mapping $T : \mathcal{X} \rightarrow \mathcal{Y}$,*

$$H(Y) \leq H(X),$$

with equality if and only if T is injective on the support of $p(x)$.

Proof. Because $Y = T(X)$ is a deterministic function of X , knowledge of X uniquely determines Y . Therefore the conditional entropy satisfies

$$H(Y \mid X) = 0.$$

By the chain rule of entropy,

$$H(X, Y) = H(X) + H(Y \mid X) = H(X).$$

But the chain rule also gives

$$H(X, Y) = H(Y) + H(X \mid Y).$$

Combining these two expressions,

$$H(X) = H(Y) + H(X \mid Y).$$

Since conditional entropy is always nonnegative, $H(X | Y) \geq 0$, it follows that

$$H(Y) \leq H(X),$$

with equality if and only if $H(X | Y) = 0$, which holds if and only if X is determined by Y —that is, if T is injective on the support of $p(x)$. $\square$

Remark 3.3. The quantity $H(X | Y) = H(X) - H(Y)$ measures the information about the source that is lost in the normalization process. It represents, in bits, the diversity of generative inputs that cannot be recovered from the normalized output. When this quantity is large, the normalization layer has compressed substantial generative diversity into a uniform representation—but the source diversity nonetheless exists and is not destroyed by the transformation.

3.4 The Data Processing Inequality

The proposition above is a special case of the more general data processing inequality, which states that no processing of a signal can increase the information it carries about its source [Shannon, 1948]. If $X \rightarrow Y \rightarrow Z$ is a Markov chain, then

$$I(X; Z) \leq I(X; Y),$$

where $I(\cdot; \cdot)$ denotes mutual information. Applied to our setting, this means that any further transformation of the normalized output Y —including reception and interpretation by a human reader—can only further reduce the information available about the generative source X .

The data processing inequality formalizes a core asymmetry in the communication system: information flows from source to receiver but cannot be increased along the way. Uniformity in Y is therefore not merely compatible with diversity in X ; it is what one should *expect* from a many-to-one normalization layer applied to a diverse source. Observing that $H(Y)$ is small provides no evidence about the magnitude of $H(X)$.

3.5 Formal Implications for the Homogenization Thesis

The implications for the homogenization thesis are direct. If LLMs operate primarily as normalization transformations applied to human-generated inputs, then the observation that LLM-assisted texts exhibit reduced stylistic variability is precisely what the channel model predicts—regardless of whether the generative source has changed at all. The reduced entropy of Y is a property of the transformation T , not of the source distribution $p(x)$.

Demonstrating cognitive homogenization requires evidence that $H(X)$ has declined, not merely that $H(Y)$ is low. These are logically independent claims. Any study that measures only properties of the outputs Y is measuring the wrong quantity.

4 Private Generative Systems and Public Registers

The formal argument of the preceding section acquires additional force when situated within the broader history of normalization in human communication systems. The coexistence of private generative diversity and public representational uniformity is not a novel feature introduced by AI systems. It is a structural characteristic of literate societies that has been reproduced across multiple historical periods and communicative technologies.

4.1 Historical Normalization Mechanisms

Spelling standardization in the early modern period dramatically reduced the orthographic diversity observable in texts without homogenizing the ideas those texts expressed. Before the consolidation of standard orthographies in the seventeenth and eighteenth centuries, written texts exhibited extraordinary variation in spelling, punctuation, and typographic practice—variation that reflected genuine diversity in regional pronunciation, scribal convention, and individual preference. The introduction of standardized dictionaries and printing practices collapsed this variation into a small set of conventional forms. Yet no historian of ideas has argued that orthographic standardization produced intellectual homogenization.

Editorial institutions in scientific publishing impose strong and consistent expectations on rhetorical structure, citation practice, and argumentative organization. A paper submitted to a journal in any major scientific discipline will pass through a normalization process that aligns its surface form with disciplinary conventions. The diversity of ideas published across such journals vastly exceeds the diversity of their rhetorical presentation. The normalization of form has not constrained the diversity of content.

Translation systems present a particularly instructive analogy. When a text is translated from one language to another, the translation necessarily loses certain features of the original—phonological texture, idiomatic resonance, lexical ambiguity—while preserving, imperfectly but substantially, the conceptual content. The translated text is a compressed projection of the original. Multiple distinct originals might produce highly similar translations. This compression does not imply that the original texts were cognitively equivalent.

4.2 LLMs as Automated Normalization Layers

Large language models can be understood as extending this historical lineage of normalization mechanisms into a new technological register. When used as editing or rewriting tools, they perform a function structurally similar to that of a human editor, style guide, or translator: they map the author’s idiosyncratic expression into a form that conforms more closely to widely shared conventions and can therefore be more easily interpreted by a broader audience.

The novelty of LLMs relative to prior normalization mechanisms lies principally in three dimensions: their computational scale, their adaptability across domains and registers, and their accessibility to users who previously lacked access to professional editorial services. None of these dimensions alters the basic logical structure of the normalization relationship between source and channel.

4.3 The Expansion of Expressive Space

A counterintuitive possibility emerges from this analysis. If large language models effectively function as translation layers between private generative systems and shared public registers, their availability may reduce the pressure that communicating agents experience to conform to normative conventions during the act of composition itself.

Prior to the availability of automated normalization, a writer who preferred an unconventional spelling system, a non-prestige dialect, or a multilingual mixing of registers faced a direct trade-off: communicate idiosyncratically and sacrifice audience, or conform to conventions and sacrifice expressive authenticity. Automated normalization dissolves this trade-off. The writer can compose in whatever register is most generatively natural and allow the normalization layer to perform the translation when communication is required.

Under these conditions, the availability of automated normalization may expand rather than contract the diversity of private linguistic practice. The apparent uniformity of the final output would reflect the existence of a translation interface, not the impoverishment of the generative source.

This possibility is not merely speculative. There is historical precedent in the introduction of word processing software, which reduced the mechanical cost of revision and thereby enabled a more exploratory approach to composition. Whether LLMs produce analogous effects on expressive diversity is an empirical question, but the theoretical possibility is well-grounded.

5 Instrumental and Epistemic Interaction

The preceding analysis focused on cases where LLMs are used instrumentally as editing or translation tools. In such cases, the human author retains primary control over the conceptual content of the communication and uses the model merely to normalize its surface form. However, LLMs are also used in modes that involve a more substantive epistemic relationship between the user and the model’s outputs. Distinguishing these modes of interaction is essential for a precise assessment of the homogenization thesis.

5.1 Instrumental Use

In instrumental use, the model functions as a tool for drafting, translation, summarization, or formatting. The user specifies the conceptual content—the ideas to be expressed, the arguments to be made, the information to be conveyed—and recruits the model to assist with the linguistic form of the expression. The model’s contribution is primarily syntactic and stylistic rather than conceptual.

Under instrumental use, the source-channel model developed above applies with maximal force. The human author’s generative processes operate largely independently of the model’s distribution, and the model’s outputs carry minimal information about those processes beyond what the author’s inputs already contain. Any convergence in the final texts reflects the normalization performed by the channel, not the convergence of the underlying sources.

5.2 Epistemic Use

In epistemic use, the user approaches the model as a source of factual information, analysis, or interpretive guidance. The user’s initial conceptual framework is incomplete or uncertain, and model-generated outputs contribute to its formation. In this mode, the model functions as a cognitive authority rather than merely a linguistic tool.

Epistemic use raises more substantive concerns about perspectival convergence. If large numbers of users consult the same systems for information and analysis, and if those systems produce consistent outputs that reflect the distributional biases of their training data, then users’ epistemic frameworks may be systematically influenced in converging directions.

However, several qualifications limit the force of this concern. First, even in epistemic use, model outputs are typically not passively accepted but are evaluated against prior beliefs, independent evidence, and alternative sources. The degree of epistemic influence depends critically on the user’s prior state of knowledge and critical orientation. Second, demonstrating lasting epistemic influence requires lon-

gitudinal evidence distinguishing temporary exposure effects from durable changes in belief or reasoning strategy—evidence that existing studies have not yet provided.

5.3 The Importance of the Distinction

The distinction between instrumental and epistemic use matters because the homogenization thesis tends to elide it. Studies that observe stylistic convergence in model-assisted texts are studying a phenomenon that occurs even in purely instrumental use, and therefore carry no implications for epistemic homogenization. Studies that observe temporary opinion alignment following exposure to opinionated model outputs are studying a phenomenon that occurs even in cases where the influence does not persist—and that therefore carry no implications for durable cognitive convergence.

A rigorous assessment of the homogenization hypothesis requires empirical designs that are sensitive to the distinction between these modes and that measure the appropriate dependent variables for each.

6 Linguistic Attractor Basins

The homogenization thesis implicitly assumes that large language models impose a single dominant attractor within the space of possible linguistic expressions: a single peak or mode in the output distribution toward which all model-assisted texts converge. This assumption is empirically false in a straightforward sense—large language models produce diverse outputs across a wide range of registers, styles, and subject matters—but the underlying structural concern deserves more precise analysis.

Definition 6.1 (Linguistic Space). Let $\mathcal{L}$ denote the space of possible linguistic expressions in a given language or family of languages. Points in $\mathcal{L}$ represent individual utterances, texts, or discourse units, endowed with whatever geometric or topological structure is appropriate to the analysis.

Definition 6.2 (Model Distribution). Let $P_\theta : \mathcal{L} \rightarrow [0, 1]$ denote the probability distribution over expressions induced by a language model with parameters θ . The model assigns higher probability to expressions that appear more frequently in configurations similar to those seen during training.

Definition 6.3 (Linguistic Attractor Basin). A *linguistic attractor basin* of a language model with distribution P_θ is a region $A \subset \mathcal{L}$ such that expressions within A have relatively high probability under P_θ and perturbations of those expressions remain within A after transformation by the normalization layer.

Several features of this definition deserve elaboration. First, attractor basins need not be unique. A language model trained on diverse stylistic corpora will develop multiple attractor regions corresponding to different registers, genres, rhetorical forms, and subject domains. The output distribution is multimodal, not concentrated at a single point.

Second, the existence of attractor basins in the model’s distribution does not imply that human generative processes collapse into those basins. When humans use language models instrumentally, their generative processes may explore large and diverse regions of $\mathcal{L}$ before the normalization layer projects the resulting expressions into nearby high-probability regions. The projection reduces entropy at the output without constraining the range of inputs.

Third, the attractor basins of a model evolve over time as the model’s parameters and fine-tuning objectives change. There is no fixed homogenizing target; the normalization surface itself is dynamic.

Proposition 6.4 (Non-Injectivity and Basin Structure). *Suppose $A \subset \mathcal{L}$ is an attractor basin of P_θ and let $U \supset A$ be a neighborhood of A in $\mathcal{L}$. If the normalization transformation T maps all points in U to points in A , then T restricted to U is non-injective for any $|U| > |A|$, and consequently*

$$H(T(X)) < H(X)$$

for any source distribution $p(x)$ supported on U .

Proof. If $|U| > |A|$ and T maps U into A , then by the pigeonhole principle, at least two distinct points in U share the same image in A under T . The transformation is therefore non-injective, and the entropy inequality $H(T(X)) < H(X)$ follows from Proposition 3.2. $\square$

7 The Conceptual Landscape

The attractor structure of a language model can be visualized geometrically by interpreting the logarithm of the model probability as a potential landscape. Define the potential function $V : \mathcal{L} \rightarrow \mathbb{R}$ by

$$V(x) = -\log P_\theta(x).$$

Regions of high probability correspond to valleys (local minima of V), while low-probability regions correspond to elevated terrain. The attractor basins identified in the previous section correspond geometrically to the valleys of this potential landscape.

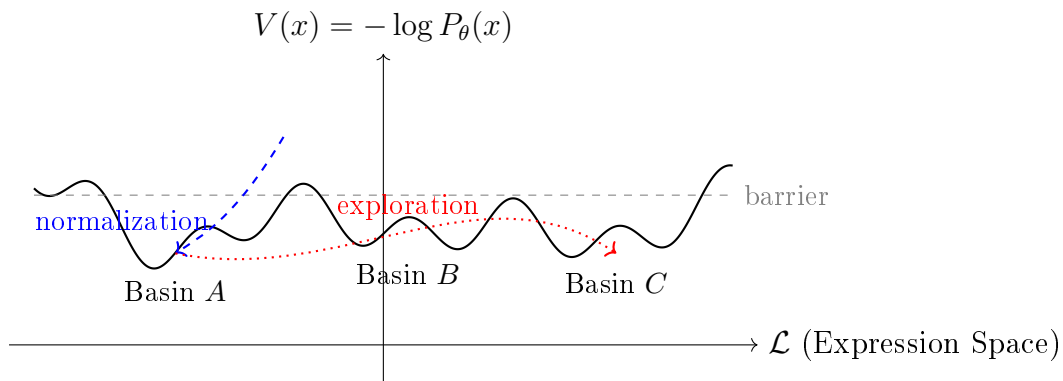


Figure 1: Potential landscape $V(x) = -\log P_\theta(x)$ showing three attractor basins A , B , C . Normalization (blue dashed) projects expressions from surrounding terrain into nearby basins. Human exploratory generation (red dotted) may traverse the landscape widely before normalization applies.

Several features of this landscape are significant for the present argument. First, models trained on diverse corpora develop landscapes with multiple valleys rather than a single global minimum. The normalized output may cluster within any of several basins rather than collapsing to a single point.

Second, human linguistic creativity characteristically involves the exploration of varied terrain: combining structures from different valleys, ascending from familiar basins to discover neighboring regions, or venturing into rarely visited high-probability subspaces that prior work has not mapped. The normalization layer applies after this exploration, projecting the resulting expression into the nearest basin without erasing the exploratory process that generated it.

Third, the landscape itself is a function of the model parameters θ , which evolve across model versions and fine-tuning objectives. The targets of normalization shift over time; there is no fixed “average” toward which all expression is inexorably drawn.

7.1 Attractor Dynamics Under Repeated Interaction

A more nuanced concern arises when users interact repeatedly with language models in a mode that is neither purely instrumental nor purely epistemic but somewhere in between. In such interactions, the model’s outputs may gradually influence the forms of expression that feel natural to the user—not by constraining the content of thought, but by providing a familiar stylistic template that the user unconsciously reproduces.

This effect is real and has plausible historical analogues. Writers who read extensively in a particular genre often find that its stylistic conventions infiltrate their own prose. Researchers trained in a specific disciplinary tradition tend to adopt its

characteristic rhetorical forms. Whether repeated interaction with model-generated text produces comparable effects is an empirical question, but the mechanism is not implausible.

The important point, however, is that such stylistic influence does not constitute cognitive homogenization in the sense relevant to the thesis. The homogenization thesis is concerned with diversity in ideas, perspectives, and reasoning strategies—the diversity of the generative source in the conceptual sense. Stylistic normalization, even if it reaches into the composition process rather than applying only as a post-hoc transformation, is compatible with the full persistence of conceptual diversity.

8 Conceptual Trajectories and Intellectual Identity

The homogenization thesis implicitly treats authorship as a property of surface linguistic style: vocabulary distributions, syntactic patterns, or rhetorical structures. Yet intellectual identity often manifests more forcefully at a deeper level—in the recurring conceptual pathways through which an author explores ideas, the persistent combinations of analytical tools they bring to bear across different problems, and the characteristic questions that organize their intellectual attention. To formalize this observation, it is useful to introduce a geometric representation of discourse in terms of trajectories through conceptual space.

8.1 Conceptual Space and the Topic Simplex

Let $\mathcal{T} = \{t_1, t_2, \dots, t_m\}$ denote the set of conceptual topics, analytical tools, and methodological frameworks available within a discourse community. Each document can be represented as a point in the probability simplex

$$\Delta^{m-1} = \left\{ (p_1, \dots, p_m) \mid p_i \geq 0, \sum_{i=1}^m p_i = 1 \right\},$$

where p_i represents the relative emphasis placed on topic t_i within the document. This representation maps each text to a vector describing its conceptual composition. A document that combines information theory, cultural evolution, and communication systems, for example, would be located in the region of Δ^{m-1} where the corresponding coordinates are jointly elevated. A document focused exclusively on a single technical apparatus would be located near a vertex.

The simplex representation abstracts away from the specific linguistic forms in which ideas are expressed. Two documents that discuss entropy and knowledge ecosystems in very different prose styles will occupy similar positions in Δ^{m-1} .

This is precisely the separation that the present analysis seeks to formalize: surface linguistic form and conceptual content can vary independently.

8.2 Conceptual Trajectories

Definition 8.1 (Conceptual Trajectory). Let $\{x_1, x_2, \dots, x_n\}$ denote the sequence of texts produced by an author over time, and let $v_k \in \Delta^{m-1}$ be the topic vector corresponding to document x_k . The sequence

$$\gamma = (v_1, v_2, \dots, v_n)$$

defines the author’s *conceptual trajectory* through the topic space.

Conceptual trajectories capture how an author’s intellectual attention moves across and between topics over time. Two authors may produce stylistically similar texts—texts that occupy nearby positions in the linguistic space $\mathcal{L}$ —while exhibiting entirely different trajectories in Δ^{m-1} . Conversely, two authors with very different stylistic signatures may trace similar paths through conceptual space.

This independence of trajectory from style is fundamental. The normalization layer applied by a language model operates on positions in $\mathcal{L}$; it does not directly alter the sequence $(v_1, v_2, \dots, v_n)$ of conceptual positions. An author’s trajectory through topic space is determined by the problems they choose to engage, the analogies they develop, and the fields they draw upon—none of which are directly constrained by a linguistic rewriting transformation.

8.3 Trajectory Distance and Author Divergence

The divergence between two authors’ intellectual practices can be quantified by a distance measure on their trajectories. One natural measure uses the average Jensen–Shannon divergence between their topic distributions over a shared window of documents:

$$D(i, j) = \frac{1}{n} \sum_{k=1}^n \text{JSD}(v_k^{(i)} \parallel v_k^{(j)}),$$

where the Jensen–Shannon divergence is defined as

$$\text{JSD}(p \parallel q) = \frac{1}{2} D_{\text{KL}}(p \parallel m) + \frac{1}{2} D_{\text{KL}}(q \parallel m), \quad m = \frac{1}{2}(p + q).$$

Unlike the Kullback–Leibler divergence, the Jensen–Shannon divergence is symmetric and bounded, making it better suited for comparing distributions that may have non-overlapping supports.

If $D(i, j)$ remains large across documents, the two authors maintain distinct conceptual identities even if stylistic features converge under normalization. The

linguistic channel compresses toward a shared register while the conceptual trajectories diverge.

8.4 Topic Constellations

Beyond individual positions in the trajectory, authors typically exhibit persistent clusters of topics that co-occur with high frequency across their bodies of work. These clusters form what may be called *topic constellations*.

Definition 8.2 (Topic Constellation). A topic constellation is a subset $C \subset \mathcal{T}$ such that topics in C co-occur with unusually high probability across the documents of a given author, relative to their marginal frequencies in the wider corpus.

Topic constellations capture stable intellectual toolkits. An author working at the intersection of information theory and cultural evolution may repeatedly combine entropy, communication channels, and knowledge ecosystems across many different texts. Another author may combine legal theory, economics, and institutional analysis. These recurring combinations are not stylistic features; they are structural features of the intellectual domain the author has carved out for themselves.

Constellations remain stable even when the surface expression of the associated texts is normalized, because the selection of conceptual topics is controlled by the author rather than the editing system.

8.5 Scaling of Conceptual Signatures

The identifiability of conceptual trajectories increases with the number of observed documents. Let $\hat{P}_i^{(n)}(t)$ denote the empirical topic distribution estimated from n documents by author i . Under standard assumptions of independent topic sampling, the law of large numbers gives

$$\hat{P}_i^{(n)}(t) \xrightarrow{n \rightarrow \infty} P_i(t) \quad \text{almost surely,}$$

and the variance of the empirical estimate satisfies

$$\text{Var}[\hat{P}_i^{(n)}(t)] = \frac{P_i(t)(1 - P_i(t))}{n} \propto \frac{1}{n}.$$

Consequently, larger bodies of work allow increasingly precise estimation of an author’s conceptual signature. This implies a counterintuitive consequence: large volumes of normalized text may strengthen, rather than weaken, the identifiability of intellectual identity. Even if stylistic features converge under model-assisted editing, the trajectory of conceptual attention becomes more precisely estimable as the number of documents increases.

Proposition 8.3 (Convergence of Empirical Topic Distributions). *Let $\hat{P}_i^{(n)}$ and $\hat{P}_j^{(n)}$ be the empirical topic distributions of authors i and j estimated from n documents each. If the true distributions P_i and P_j are distinct, then for any $\epsilon > 0$ there exists N such that for all $n > N$,*

$$\mathbb{P}\left(\text{JSD}(\hat{P}_i^{(n)} \parallel \hat{P}_j^{(n)}) > 0\right) > 1 - \epsilon.$$

This proposition formalizes the claim that distinct conceptual identities remain distinguishable in the limit of large text corpora, regardless of surface stylistic normalization.

9 Topic Monotonicity and Conceptual Toolchains

The trajectory framework of the previous section captures how authors move through conceptual space over time. However, intellectual practice exhibits a further structural regularity: authors not only visit characteristic conceptual regions; they recurrently deploy the same sequences of analytical operations across different problems. This recurring structure can be described as *topic monotonicity*.

9.1 Conceptual Toolchains

Definition 9.1 (Conceptual Toolchain). A *conceptual toolchain* is an ordered sequence

$$\tau = (t_{i_1}, t_{i_2}, \dots, t_{i_k})$$

of conceptual tools drawn from $\mathcal{T}$ that are applied in a recognizable logical order during an argument or analysis.

Conceptual toolchains represent stable intellectual procedures. In information-theoretic analyses of social phenomena, a characteristic sequence might proceed: probability distribution $\rightarrow$ entropy measure $\rightarrow$ channel model $\rightarrow$ inequality derivation $\rightarrow$ cultural interpretation. In economic analysis of institutions, a recurring chain might move: incentive structure $\rightarrow$ equilibrium concept $\rightarrow$ comparative statics $\rightarrow$ welfare evaluation. In historical linguistics, a standard sequence might combine phonological reconstruction $\rightarrow$ morphological analysis $\rightarrow$ contact hypothesis $\rightarrow$ grammatical typology.

These toolchains are not merely stylistic habits. They represent structured analytical methods that authors have developed, refined, and repeatedly applied. Their persistence across texts is evidence of genuine intellectual commitment to a set of procedures—a deeper form of identity than any stylistic signature.

9.2 Topic Monotonicity

Definition 9.2 (Topic Monotonicity). An author’s work exhibits *topic monotonicity with respect to toolchain* τ if

$$\mathbb{P}(\tau \mid i) \gg \mathbb{P}(\tau),$$

where $\mathbb{P}(\tau \mid i)$ is the probability that author i employs toolchain τ in a given document and $\mathbb{P}(\tau)$ is the background frequency of τ across the wider corpus.

When an author exhibits strong topic monotonicity, their characteristic toolchains function as intellectual fingerprints—analogous to recurring motifs in musical composition or consistent compositional choices in visual art: elements that reveal the systematic structure of a creative practice without exhausting its content.

The detectability of topic monotonicity is a direct consequence of the convergence result in Proposition 8.3. As the number of documents increases, the empirical frequency of each toolchain stabilizes, and differences between authors become increasingly apparent.

9.3 Normalization and Toolchain Preservation

Normalization transformations applied by language models operate primarily on surface linguistic structure. They affect lexical selection, sentence structure, and rhetorical organization, but they typically leave intact the underlying logical and methodological sequence of an argument. The transformation $T : X \rightarrow Y$ introduced in Section 3 can therefore be understood as a projection that preserves the sequence of conceptual operations embedded in the input even as it rewrites the linguistic form of their expression.

This claim rests on the observation that conceptual toolchains are encoded at a level of abstraction above the lexical and syntactic features that normalization modifies. The step from “probability distribution” to “entropy measure” is a logical relationship between conceptual objects. Whether the associated sentence reads “the entropy of the source distribution” or “the information content of the generative process” does not alter the conceptual operation being performed. Both phrasings can arise from the same underlying analytical step; the normalization layer may choose between them without changing the structure of the argument.

This asymmetry implies that stylistic similarity among normalized texts is a weak signal about the conceptual procedures that generated them. Two texts that employ the same conceptual toolchain may indeed converge stylistically after normalization. But the converse does not hold: stylistically similar outputs may arise

from entirely different analytical sequences. The homogenization of form does not imply the homogenization of procedure.

9.4 Intellectual Identity as a Structured Conceptual Toolkit

The combined structure of conceptual trajectories, topic constellations, and topic monotonicity suggests a more comprehensive account of intellectual identity than models based solely on stylistic analysis. Within this framework, authorship is not defined primarily by the surface characteristics of linguistic expression but by the recurrent configuration of conceptual tools through which an individual engages with problems.

An author’s intellectual identity can be understood as a structured conceptual toolkit that evolves through repeated engagement across many texts. This toolkit includes the particular analytical tools the author tends to employ, the characteristic combinations through which those tools are brought into relation, and the recurring sequences of conceptual operations that structure argument development. Over time, these patterns generate stable intellectual signatures that persist across bodies of work even when the linguistic surface of individual documents is partially standardized.

Within this perspective, stylistic normalization affects the representational surface of discourse without erasing the deeper organizational patterns that characterize intellectual practice. The persistence of these patterns indicates that intellectual identity is encoded primarily in the structure of conceptual activity rather than in the specific linguistic forms through which that activity is expressed.

10 Multimodal Signatures and Input-Layer Diversity

The homogenization thesis typically treats the textual output of a communication system as the primary observable signal of authorship. Yet communication is produced through embodied processes—speech, handwriting, typing, gesture, and interaction timing—each of which leaves distinctive behavioral traces that are logically independent of the linguistic representation ultimately transmitted. These input modalities introduce an additional dimension of diversity that remains largely unaffected by the normalization layers applied to textual output.

10.1 Behavioral Signals in Communication

Let B denote the set of behavioral signals associated with the production of a communicative act. These signals arise from the motor and perceptual processes

through which linguistic and conceptual content is generated and transmitted. They include the acoustic characteristics of speech production, the geometric and kinematic trajectories of handwriting, the temporal structure of keystroke dynamics during typing, patterns of gestural emphasis accompanying explanation, and the rhythmic interaction traces that accumulate during composition and revision.

Each class of behavioral signal captures a distinct aspect of the underlying generative process. Speech production contains stable acoustic features associated with the anatomical configuration of the vocal tract and with habitual patterns of articulation learned over a lifetime. These features are sufficiently distinctive that speaker identification from short voice samples achieves high accuracy even under varied recording conditions. Handwriting trajectories encode both the geometric shape of letterforms and the dynamic movement properties—velocity, pressure, angle—that reflect learned motor patterns. Typing behavior exhibits consistent temporal structures in interkeystroke intervals and pause distributions that remain stable for a given typist across different texts and sessions. Gestural interaction traces encode timing and spatial patterns of attention that correlate with cognitive processing styles.

The statistical distinctiveness of these signals has been extensively documented in the biometric identification literature. The relevant finding for the present argument is that behavioral signals originate in processes that are substantially more resistant to normalization than the linguistic representations they help to produce.

10.2 Generative Model of Expression

The production of a textual expression can be represented as a joint generative process

$$X = G(C, B),$$

where $C \in \mathcal{C}$ denotes the conceptual content generated by the author’s cognitive processes and B denotes the behavioral signals associated with the physical production of the expression. The function G maps the joint state (C, B) to a linguistic expression X .

Large language models interact exclusively with the textual representation X . When used as normalization layers, they apply a transformation $Y = T(X)$ that modifies the linguistic surface while leaving the upstream variables C and B unchanged. The behavioral signals B are invisible to the model; the conceptual content C passes through the model’s rewriting without being directly reconstructed or constrained.

The entropy of the full generative system is therefore

$$H(C, B) = H(C) + H(B \mid C) \geq H(C),$$

and the normalization layer reduces only the entropy of the observable representation:

$$H(Y) = H(T(X)) \leq H(X) \leq H(C, B).$$

The gap between $H(C, B)$ and $H(Y)$ measures the diversity that is compressed by the normalization process but that remains present in the upstream generative system.

10.3 Multimodal Authorship Signatures

Communication environments increasingly support the simultaneous use of multiple input modalities. A written document may appear alongside a spoken explanation, handwritten annotations produced during drafting, and interaction traces recording the temporal structure of the composition process. Together, these heterogeneous signals form a high-dimensional representation of authorship. Each modality captures a distinct aspect of the expressive process; when considered jointly, they provide a far richer description of communicative activity than any single channel alone.

The combination of modalities produces what may be described as a high-dimensional authorship signature. Even if individual modalities undergo partial normalization, the joint distribution of signals across textual, acoustic, motor, and temporal channels retains substantial entropy. The probability that two individuals would independently produce identical patterns across all of these channels simultaneously is negligibly small.

Consequently, multimodal communication systems may enable new forms of authorship attribution that extend beyond traditional stylometric analysis. Rather than relying solely on surface linguistic features, such systems incorporate behavioral and temporal signals that originate in the embodied processes of communication. These signals provide additional structure through which individual patterns persist even in environments where linguistic surfaces become increasingly standardized.

10.4 Scaling Diversity Through Input Channels

The presence of multiple input modalities implies that the effective dimensionality of the expressive space grows with the number of channels available to a communicator. Each additional modality introduces new behavioral degrees of freedom into the generative process. The entropy of the full multimodal system is bounded below by the entropy of any single channel and above by the sum of marginal entropies:

$$\max_b H(b) \leq H(B) \leq \sum_b H(b_k).$$

Large language models compress only one projection of this high-dimensional system: the textual representation transmitted through the written channel. The broader generative system—comprising conceptual trajectories, multimodal behavioral signals, and temporal production patterns—remains high-dimensional and richly varied.

11 Dialect Liberation Under Automated Normalization

A central assumption of the homogenization thesis is that the widespread use of large language models will encourage individuals to converge toward a single dominant linguistic register. This assumption rests on an implicit model of communication in which speakers must adopt shared conventions to achieve intelligibility across heterogeneous audiences. It overlooks a structural change introduced by automated normalization: when translation and rewriting layers are available, the communicative cost of producing nonstandard language decreases substantially, and this reduction may expand rather than contract the diversity of upstream linguistic expression.

11.1 Communication Friction and Standardization Pressure

Linguistic standardization has historically been driven by communication friction: the interpretive cost borne by listeners when they encounter unfamiliar forms. If the listener cannot decode an unconventional spelling, dialectal construction, or mixed-register expression, communication fails. Under such conditions, speakers face a practical incentive to adopt conventions that minimize this friction, even at the cost of suppressing their own more natural or efficient linguistic habits.

Formally, let $C(x, y)$ denote the cost of interpreting expression x for listener y . The expected communication cost over a population of listeners is

$$E[C] = \sum_{x,y} P(x, y) C(x, y).$$

Standardization reduces this expected cost by restricting the distribution over expressions x to forms that listeners can reliably decode. The resulting reduction in expected cost is a genuine benefit, but it is purchased at the price of compressing the distribution of upstream expression toward a smaller set of conventional forms.

11.2 Automated Normalization as a Cost-Reducing Layer

Automated normalization alters this cost structure. When a generative model can reliably rewrite an idiosyncratic expression into a widely interpretable standard form, the listener no longer bears the full burden of decoding nonstandard input. The effective interpretation cost of the original expression X is reduced to the interpretation cost of the normalized output $Y = T(X)$:

$$C'(X, y) = C(T(X), y).$$

If T reliably maps diverse inputs to forms that audiences can understand, the expected communication cost $E[C']$ decreases even when the original expressions remain highly variable.

This means that the practical incentive for speakers to conform to standardized conventions during composition weakens when normalization layers are available. The friction that historically drove standardization is decoupled from the composition stage and relocated to the output stage, where it is absorbed by the automated system.

11.3 Expansion of Private Linguistic Systems

When the communicative cost of nonstandard expression decreases, speakers acquire greater freedom to produce language in forms that reflect their own linguistic habits, local environments, or idiosyncratic cognitive preferences. Such private linguistic systems can take many forms: regional dialects reflecting local speech communities, hybrid multilingual constructions blending two or more languages within a single drafting environment, unconventional spelling conventions adopted for efficiency or habit, personal shorthand developed through repeated writing practice, or partially formed provisional expressions that function as cognitive scaffolding during early drafting stages.

The introduction of word processing software offers a useful historical parallel. By reducing the mechanical cost of revision, word processors encouraged more exploratory drafting behaviors: writers could generate provisional expressions rapidly and refine them iteratively, rather than composing polished sentences at the outset. Automated linguistic normalization extends this pattern further, allowing authors to experiment with unconventional linguistic structures during composition while relying on the normalization layer to convert the result into a standard communicative register when needed.

11.4 Decoupling of Generative and Communicative Registers

The communication structure that emerges under automated normalization introduces a separation between two layers of linguistic activity that were historically more tightly coupled. The *private generative register* corresponds to the linguistic system employed by the author during composition: provisional, locally adapted, potentially idiosyncratic, and not directly constrained by audience legibility requirements. The *public communicative register* corresponds to the representation ultimately transmitted to readers, which is required to satisfy broadly shared intelligibility standards.

Historically, the need to produce intelligible documents required authors to convert their private register into standardized forms before publication. This conversion imposed costs that encouraged writers to adopt conventional forms earlier in the composition process, effectively compressing the generative stage toward the communicative standard.

Large language models dissolve this compression by performing the conversion automatically. The public representation may appear stylistically standardized while the generative register remains heterogeneous and locally adapted. The apparent convergence of linguistic outputs within model-mediated systems may therefore conceal rather than reflect the structure of the generative processes that produced them.

Proposition 11.1 (Decoupling of Registers Under Normalization). *Let X be a random variable representing expressions drawn from a private generative register with entropy $H(X)$, and let $Y = T(X)$ be the normalized public representation. The correlation between $H(X)$ and $H(Y)$ over a population of authors is not constrained to be positive. In particular, there exist systems in which $H(Y)$ decreases while $H(X)$ remains constant or increases.*

Proof. By Proposition 3.2, $H(Y) \leq H(X)$ for any individual author. If a normalization technology reduces the communicative cost of private diversity, it may encourage authors to increase the entropy of X while the entropy reduction imposed by T ensures that $H(Y)$ remains low. These two effects operate on different variables and can vary independently. $\square$

11.5 Implications for Linguistic Diversity

If automated normalization continues to improve, the incentive for authors to conform to prestige dialects during composition may decline further. Individuals may increasingly compose in the registers most cognitively natural or locally appropriate

for them, relying on normalization systems to translate those registers into shared communicative forms when necessary.

The systems that appear to standardize linguistic output may therefore simultaneously contribute to the diversification of upstream generative practice. The observable uniformity of published texts would in this scenario reflect the increasing sophistication of the translation layer, not the impoverishment of the generative diversity that feeds it.

12 Constraint Geometry of Conceptual Spaces

The preceding sections have analyzed linguistic and conceptual diversity primarily through the lens of probability distributions and entropy. A complementary perspective emerges from interpreting the space of possible expressions as a geometric structure shaped by constraints. This geometric viewpoint illuminates a further reason why surface normalization cannot be taken as evidence of deep conceptual convergence.

Let $\mathcal{L}$ denote the space of linguistic expressions and $\mathcal{C}$ the space of conceptual configurations available within a discourse community. Rather than treating these spaces as unstructured sets, we may view them as manifolds whose geometry is shaped by constraints arising from grammatical well-formedness, shared background knowledge, disciplinary conventions, and cognitive limitations on complexity.

Under this interpretation, generative processes do not explore the ambient spaces uniformly. Instead, they trace paths along structured submanifolds corresponding to coherent conceptual frameworks or intellectual traditions. Each tradition defines a locally valid coordinate system within a region of the conceptual manifold; related ideas cluster along shared geometric directions, while ideas from distinct traditions inhabit different regions of the manifold.

12.1 Normalization as Projection onto the Linguistic Manifold

Normalization processes performed by language models can be understood as projections onto high-probability regions of the linguistic manifold $\mathcal{L}$. Such projections alter the position of an expression within $\mathcal{L}$ without necessarily altering its conceptual location in $\mathcal{C}$.

Definition 12.1 (Normalization Projection). Let $\pi : \mathcal{L} \rightarrow \mathcal{L}_{\text{norm}}$ denote the projection of a linguistic expression onto the high-probability submanifold $\mathcal{L}_{\text{norm}} \subset \mathcal{L}$ defined by the language model’s output distribution. A normalization transformation is a map of the form $T = \pi \circ \phi$ where $\phi : \mathcal{C} \rightarrow \mathcal{L}$ encodes the conceptual content in linguistic form.

The key geometric observation is that the projection π can reduce the variance of observed expressions within $\mathcal{L}$ while leaving the distribution over $\mathcal{C}$ unchanged. Multiple conceptual inputs, mapped to different points in $\mathcal{L}$ by ϕ , may be projected by π to nearby points in $\mathcal{L}_{\text{norm}}$.

12.2 Strata and Cross-Stratum Transition

The conceptual manifold $\mathcal{C}$ exhibits a stratified structure in which different intellectual traditions occupy partially overlapping submanifolds, called strata. Within any given stratum, conceptual trajectories tend to follow directions that are locally consistent with existing frameworks: refining definitions, extending models, applying established theories to new domains.

Innovation frequently arises when a trajectory departs from these locally dominant directions and moves transversely between strata. Such transverse movements correspond to conceptual recombination across previously separated frameworks. The introduction of information-theoretic tools into biology, the application of evolutionary reasoning to cultural phenomena, or the incorporation of category-theoretic structures into programming language semantics all represent such transverse transitions.

Stylistic convergence within $\mathcal{L}$ is compatible with the full preservation of cross-stratum movement in $\mathcal{C}$. The projection π may map diverse conceptual inputs to similar linguistic forms without collapsing the stratified structure of the conceptual manifold. Intellectual innovation—which depends on transverse movement between strata—is not constrained by the uniformity of the linguistic representation that describes it.

13 Expressive Degeneracy in Linguistic Systems

The sections above have emphasized the logical independence of generative diversity from representational uniformity. A structural explanation for this independence can be given in terms of *expressive degeneracy*: the phenomenon whereby many distinct conceptual configurations produce the same or similar observable linguistic representations. Degeneracy is widespread in biological, physical, and computational systems; in the context of language, it explains why stylistic convergence at the surface level cannot be taken as evidence of convergence in the underlying generative processes.

13.1 Many-to-One Structure of Expressive Mapping

Let $\mathcal{C}$ denote the space of conceptual configurations and $\mathcal{L}$ the space of linguistic expressions. The act of writing can be formalized as a mapping

$$F : \mathcal{C} \rightarrow \mathcal{L}$$

that translates conceptual structures into linguistic form. This mapping is neither injective nor uniquely determined by C : multiple distinct conceptual configurations can give rise to the same linguistic expression (many-to-one), and a single conceptual configuration can typically be expressed through multiple linguistic forms (one-to-many). The mapping F is therefore highly non-injective in both directions.

When a normalization layer T is applied to produce the observable representation $Y = T(F(C))$, the degree of non-injectivity compounds further. Distinct conceptual inputs $C_1, C_2, \dots$ may produce distinct expressions $F(C_1), F(C_2), \dots$ subsequently projected by T onto the same normalized output Y . The equivalence class of conceptual configurations producing a given normalized output is

$$[Y] = \{C \in \mathcal{C} : T(F(C)) = Y\},$$

and this class may be large.

Proposition 13.1 (Degeneracy Under Composed Normalization). *If F is many-to-one and T is many-to-one, then the composed map $T \circ F$ is many-to-one and $|[Y]| \geq \max(|F^{-1}(T^{-1}(Y))|)$. The degeneracy of the full system is at least as large as the degeneracy of either component alone.*

The proof follows immediately from the definition of composition: distinct inputs that map to the same intermediate expression will certainly map to the same output, and distinct inputs that map to different intermediate expressions may still map to the same output if T collapses those expressions.

13.2 Implications for Interpretability of Output Similarity

From the perspective of an observer who sees only the normalized output Y , the equivalence class $[Y]$ is entirely opaque. Observing two texts with the same normalized representation provides no information about whether they belong to the same equivalence class for conceptual reasons (similar ideas expressed similarly) or as a consequence of degeneracy (diverse ideas that happen to produce the same normalized form).

A useful analogy is the compilation of software source code into machine instructions. Distinct implementations of the same algorithm—or of entirely different

algorithms that happen to produce the same output on a given input—may compile to nearly identical sequences of machine instructions once comments, variable names, and higher-level structural differences are eliminated. The similarity of the compiled form reveals little about the diversity of the original programs. Linguistic normalization operates analogously: it compresses diverse generative structures into common representational forms without eliminating the diversity of the structures themselves.

13.3 Degeneracy Grows with Dimensionality

In high-dimensional conceptual spaces, degeneracy increases rapidly with dimension. As the number of independent conceptual dimensions m increases, the number of distinct configurations that project to any given low-dimensional representation grows. The compression introduced by normalization therefore conceals an expanding rather than contracting diversity of generative possibilities.

This observation is relevant to the specific context of contemporary LLMs. These systems produce text from a high-dimensional internal state space but output a relatively low-entropy sequence of tokens. The resulting compression from internal state space to textual output introduces substantial degeneracy even before the human author’s conceptual diversity is considered. The addition of the human generative layer further amplifies the size of the equivalence classes associated with any given normalized output.

Recognizing the role of expressive degeneracy clarifies the structure of the homogenization debate. The question is not whether linguistic representations can become standardized—they evidently can—but whether the generative diversity of the conceptual space is itself contracting. The presence of large degeneracy classes in the mapping from conceptual to linguistic space shows that these two phenomena can vary entirely independently.

14 Combinatorial Expansion in Conceptual Space

Expressive degeneracy explains why a compressed linguistic representation is compatible with diverse underlying conceptual configurations. A related but distinct argument concerns the sheer combinatorial scale of the conceptual space itself: even if normalization systems were perfectly transparent—if every distinct conceptual configuration produced a distinct observable output—the space of available configurations would remain astronomically large and would not be meaningfully constrained by the existence of the normalization layer.

14.1 Combinatorial Structure of Topic Combinations

Let $\mathcal{T} = \{t_1, t_2, \dots, t_m\}$ denote the set of conceptual topics available within an intellectual community. Conceptual configurations can be represented as subsets or structured combinations of elements from $\mathcal{T}$. Even restricting attention to unordered subsets, the number of possible configurations is

$$|\mathcal{P}(\mathcal{T})| = 2^m.$$

For ordered sequences of length k , the number of possible conceptual toolchains is m^k . In practice, m is large—the relevant conceptual space of a modern intellectual community encompasses thousands of distinct topics, methods, and frameworks—so the combinatorial scale of the configuration space is effectively unlimited.

The key point is that the normalization layer does not reduce the size of this conceptual space. It acts on the linguistic representations of configurations, not on the configurations themselves. The 2^m possible conceptual subsets remain available to any author; the normalization layer merely reshapes how the resulting expressions are rendered in standard form.

14.2 Coordination of Representation and Expansion of Content

The parallel with mathematical notation is instructive. Standard mathematical notation imposes strict conventions on symbolic expression: specific symbols for common operations, conventional reading order, required typographic formats. These conventions dramatically reduce the stylistic variation observable in mathematical texts. Yet the conceptual content expressible within this standardized notation is not similarly constrained. New theorems, novel proof strategies, and unexpected mathematical structures are discovered continuously within the same representational system that was in use when those structures were unknown.

The standardization of notation coordinates expression without constraining content. This is possible precisely because the space of conceptual configurations is vastly larger than the space of representational forms. Linguistic normalization by LLMs plays an analogous role: it standardizes the representational layer while leaving the conceptual layer free to expand.

14.3 Innovation from Rare Combinations

Intellectual innovation characteristically arises from combinations of ideas that were previously rare or unexplored. New research programs, theoretical frameworks, and interdisciplinary syntheses emerge when conceptual tools from different regions of

$\mathcal{T}$ are combined in configurations that have not previously been attempted.

Proposition 14.1 (Expansion of Innovative Configurations). *Let $\mathcal{T}$ be a set of m conceptual topics. Define an innovative configuration as any subset of $\mathcal{T}$ that has not previously been observed in the corpus. The number of innovative configurations grows at least as fast as $2^m - |\mathcal{C}_{\text{obs}}|$, where $|\mathcal{C}_{\text{obs}}|$ is the number of previously observed configurations. For fixed $|\mathcal{C}_{\text{obs}}|$, this quantity grows exponentially in m .*

The proposition implies that as the conceptual vocabulary available to a community grows, the number of unexplored combinations grows even faster. Intellectual diversity in this sense is structurally self-amplifying, not self-limiting. The introduction of linguistic normalization does not constrain this combinatorial expansion because the normalization operates on the representational surface rather than on the conceptual vocabulary.

15 Transverse Innovation in Conceptual Manifolds

The combinatorial argument of the previous section establishes that the potential space of conceptual configurations is effectively unbounded. The present section complements this with a dynamical argument: within actual intellectual practice, innovation arises through specific kinds of movement across the conceptual landscape that are not constrained by linguistic normalization.

15.1 Longitudinal versus Transverse Intellectual Movement

Within any given conceptual stratum—any established tradition, disciplinary framework, or theoretical program—intellectual work proceeds primarily in *longitudinal directions*: extending existing models, clarifying their implications, testing them against new evidence, and resolving internal inconsistencies. This kind of work refines and develops an intellectual tradition from within.

Innovation of a different character arises through *transverse movement*: the departure from one established region of the conceptual manifold and the entry into a neighboring region, establishing connections between previously separated frameworks. Such transverse transitions are the mechanism by which new intellectual traditions come into existence.

Historical examples are abundant. The mathematical formalization of probability in the seventeenth century required transverse movement between geometry, combinatorics, and practical reasoning about chance. The development of evolutionary epistemology required transverse movement between biology and philosophy of science. The emergence of cognitive neuroscience required transverse movement between experimental psychology, computational modeling, and neuroanatomy. In

each case, the originality consisted not in moving further along an established path but in crossing from one region of the conceptual landscape into a neighboring one.

15.2 Independence of Transverse Movement from Stylistic Form

Transverse innovation operates in the conceptual space $\mathcal{C}$ and is defined by the structure of the transition between strata. The linguistic representation of such a transition may be expressed through highly standardized prose without losing its novelty. The originality of an interdisciplinary synthesis lies in the conceptual linkage it establishes, not in the rhetorical form of its expression.

This observation further weakens the inference from stylistic convergence to cognitive homogenization. Even if normalization processes encourage similar rhetorical structures across documents, the topology of intellectual exploration—including the occurrence of transverse movements between strata—may remain highly diverse. Two documents may appear stylistically similar after normalization while one represents a transverse conceptual move of significant originality and the other represents a routine longitudinal extension.

15.3 Preservation of Transverse Directions

The availability of automated normalization does not reduce the number of transverse directions available in the conceptual manifold. The manifold’s structure is determined by the intellectual traditions and frameworks present in the discourse community, not by the linguistic register through which those traditions are expressed. As new disciplines form, new conceptual strata emerge, and new transverse directions between strata become available for exploration.

Moreover, the reduction of linguistic friction may facilitate transverse movements by making it easier for scholars from different traditions to communicate across disciplinary boundaries. When an information theorist and a cultural anthropologist can express their respective frameworks in mutually legible registers with lower translation costs, the conditions for productive cross-stratum recombination improve.

16 Exploratory Entropy and the Ecology of Knowledge

The preceding analysis has argued that generative diversity at the level of the cognitive source is unlikely to be substantially reduced by channel normalization. This does not, however, exhaust the culturally relevant consequences of LLM-mediated

communication. A distinct and more subtle form of homogenization may operate not at the level of expression but at the level of *discovery*: the pathways through which individuals encounter ideas, encounter conflicting interpretations, and develop the conditions for genuine conceptual innovation.

16.1 Exploratory Entropy Defined

Definition 16.1 (Knowledge Environment). Let $\mathcal{S} = \{s_1, s_2, \dots, s_n\}$ denote the set of informational sources available within a knowledge environment. These include books, articles, archival materials, personal correspondence, course lectures, and any other medium through which ideas are encountered. Let $P(s_i)$ denote the probability that a learner encounters source s_i during an episode of intellectual exploration.

Definition 16.2 (Exploratory Entropy). The *exploratory entropy* of a knowledge environment is

$$H_E = - \sum_{s \in \mathcal{S}} P(s) \log P(s).$$

High exploratory entropy corresponds to a discovery process in which exposure is distributed broadly across the available sources; low exploratory entropy corresponds to a process in which a small subset of sources dominates exposure, regardless of the size of $\mathcal{S}$.

Remark 16.3. Exploratory entropy is distinct from both source entropy $H(X)$ and channel entropy $H(Y)$. It concerns the diversity of informational trajectories through which ideas are encountered, rather than the diversity of expressions produced or received. A system can maintain high $H(X)$ and high $H(Y)$ while exhibiting low H_E —that is, individuals may continue to produce and receive diverse expressions while all encountering the same small set of ideas along the same conceptual pathways.

16.2 Traditional Knowledge Environments

Traditional knowledge environments—physical libraries, disciplinary reading lists, archival research, informal scholarly correspondence—characteristically maintained relatively high exploratory entropy. The distribution $P(s)$ was shaped by the intersection of multiple independent factors: geographic proximity to particular collections, disciplinary training, personal recommendations, and the productive chance of finding unexpected materials in the course of searching for something else. These factors created overlapping but partially independent pathways through the available corpus, exposing different individuals to different subsets of available ideas and thereby generating the conditions for diverse intellectual trajectories.

High exploratory entropy also enables what might be called *cross-basin discovery*: the encounter with ideas that inhabit a different region of the conceptual landscape from the region one’s prior work occupies. Such encounters are a significant source of intellectual innovation, since they introduce unexpected analogies, challenge presupposed frameworks, and open new angles of approach to persistent problems.

16.3 Algorithmic Mediation and the Concentration of Discovery

Generative AI systems, by contrast, typically synthesize information from many sources into a single summarized representation. The user who queries a language model for an explanation of a complex topic receives a single canonicalized output that collapses the diversity of sources and interpretations into one streamlined narrative.

Formally, let $\mathcal{C}$ be a corpus of documents. A generative summary $\mathcal{G}(\mathcal{C})$ is a compression operator

$$\mathcal{C} \rightarrow \mathcal{G}(\mathcal{C})$$

that maps the full corpus onto a single synthesized representation. Users who rely on $\mathcal{G}(\mathcal{C})$ rather than directly consulting the corpus encounter a highly concentrated distribution over ideas: the ideas that the synthesis surface prominently, weighted by their frequency and the model’s training distribution.

If a large proportion of users rely on the same generative systems to navigate the same knowledge domains, the distribution $P(s)$ over sources becomes increasingly concentrated around the canonical outputs of those systems. Exploratory entropy declines—not because the corpus has shrunk or because individuals have become less capable of diversity in their own expression, but because the pathways through which they encounter the corpus have contracted.

Proposition 16.4 (Entropy Concentration Under Summarization). *Let $\mathcal{C} = \{c_1, \dots, c_m\}$ be a corpus with uniform discovery distribution $P_0(c_i) = 1/m$. Suppose a generative summary assigns saliency weights $w_i \geq 0$ with $\sum_i w_i = 1$, defining the mediated distribution $P_G(c_i) = w_i$. If the weights are nonuniform, i.e. $\max_i w_i > 1/m$, then*

$$H_E(P_G) < H_E(P_0) = \log m.$$

Proof. Since entropy is uniquely maximized by the uniform distribution over a finite set, any nonuniform distribution P_G over $\{c_1, \dots, c_m\}$ satisfies $H_E(P_G) < \log m = H_E(P_0)$. $\square$

16.4 The Significance of Exploratory Variance

The cultural significance of exploratory entropy lies in its relationship to intellectual innovation and adaptive resilience. From the perspective of cultural evolution, the resilience of a knowledge ecosystem depends not only on the diversity of ideas it currently contains but on the probability that novel combinations and reframings will be generated in the future. This probability is higher when exploratory pathways are diverse, because diversity in discovery creates the conditions for unexpected cross-fertilization.

If the mechanisms of discovery increasingly converge toward centralized systems that surface the same canonical framings to large populations of users, the expected rate of conceptual recombination may decline. Many users may encounter similar interpretive schemas even if their individual expressive styles remain entirely distinct. The cultural consequence would not be a reduction in stylistic variation but a narrowing of the conceptual routes through which ideas are encountered—a form of homogenization that operates at the level of the epistemic infrastructure rather than at the level of individual expression.

17 Rank-Mediated Discovery and Its Entropy

Modern information environments rarely expose users directly to the full set of available sources. Instead, search engines, recommender systems, and generative models mediate access through ranking, filtering, and summarization mechanisms. Each of these mechanisms systematically transforms the discovery distribution in ways that reduce exploratory entropy.

17.1 Ranking Functions and Exposure Distributions

Let $\mathcal{S} = \{s_1, \dots, s_n\}$ be a set of available sources. A ranking function R assigns to each source an ordinal rank $R(s_i) \in \{1, 2, \dots, n\}$. A monotonically decreasing exposure function $f : \{1, \dots, n\} \rightarrow (0, 1]$ translates ranks into exposure probabilities. The rank-mediated discovery distribution is

$$P'(s_i) = \frac{f(R(s_i))}{\sum_{j=1}^n f(R(s_j))}.$$

In practice, f decays sharply with rank position. Studies of user behavior on search engines consistently find that the probability of clicking on a result falls off rapidly beyond the first few ranked items [Lazer et al., 2009, Raghavan and Hanna, 2023]. A first-ranked source may receive orders of magnitude more exposure than a tenth-ranked source.

Proposition 17.1 (Ranking Reduces Exploratory Entropy). *Let P_0 be the uniform distribution over $\mathcal{S}$. For any monotonically decreasing exposure function f that is non-constant,*

$$H_E(P') < H_E(P_0) = \log n.$$

Proof. The rank-mediated distribution P' is nonuniform whenever f is non-constant, since sources at different rank positions receive different exposure probabilities. Entropy is uniquely maximized by the uniform distribution, so $H_E(P') < \log n = H_E(P_0)$. $\square$

17.2 Compounding Effects of Ranking and Summarization

Generative AI systems introduce an additional layer of compression beyond ranking. Rather than merely ordering sources, they synthesize multiple ranked sources into a single output. If the user receives only the synthesized output rather than the ranked list, the effective discovery distribution collapses from P' to a point mass on the synthesized summary—an entropy of zero.

In practice, users retain some residual exposure to original sources through citations and links, and some proportion of users actively pursue original materials. The effective discovery distribution is a mixture between the point-mass on the summary and the original distribution P' over cited sources. But even in the most optimistic case, the entropy of this mixture is bounded above by the entropy of P' , which is itself bounded by $\log n$. Summarization introduces a floor on entropy reduction that stacks multiplicatively with the reduction introduced by ranking.

17.3 Design as a Variable

It is important to note that the entropy of the rank-mediated discovery distribution depends on the specific design of the information interface, not merely on the existence of algorithmic mediation. Interfaces that surface multiple alternative summaries, provide explicit citation links to diverse original sources, incorporate stochastic exploration mechanisms, or present conflicting interpretations alongside consensus views will maintain higher exploratory entropy than interfaces that present a single canonical response.

The policy implications are therefore not that generative AI systems necessarily reduce exploratory entropy, but that their design choices determine by how much they reduce it. This is a tractable engineering and institutional problem, not a structural inevitability.

18 Normalization Layers in Historical Communication Systems

The analysis developed in this paper situates large language models within a long historical tradition of normalization mechanisms in communication systems. This situating serves a dual function: it both demystifies the LLM case by showing that its structural logic is not novel, and it furnishes historical evidence relevant to the empirical question of whether normalization mechanisms have historically produced cognitive convergence.

18.1 Case Studies in Communicative Standardization

Three historical cases merit detailed consideration.

Standardized orthography. The emergence of standardized spelling in European languages between the sixteenth and eighteenth centuries compressed an extraordinary range of orthographic variation into a small set of conventional forms. The linguistic diversity of the pre-standardization period—in which educated writers might spell the same word differently on the same page—is well documented by historians of language. Its compression did not visibly impoverish the diversity of ideas expressed in the resulting standardized texts. The century following orthographic standardization in England was the century of the Scientific Revolution. If anything, standardization facilitated the cross-regional communication of diverse ideas by reducing the cognitive cost of decoding nonstandard orthographic forms.

Scientific editorial conventions. The development of standardized scientific prose in the late seventeenth and eighteenth centuries—associated with institutions such as the Royal Society and the Académie des sciences—imposed strong rhetorical norms on scientific writing. The passive voice, the impersonal register, the systematic use of tables and mathematical notation, and the conventional structure of hypothesis, method, result, and discussion are all artifacts of this normalization process. None of these conventions has been shown to have constrained the diversity of scientific ideas expressed within them; on the contrary, they facilitated the communication of diverse findings across linguistic and disciplinary barriers.

Translation. The practice of translation has always involved the compression of generative diversity. A translation system—whether human or automated—maps the rich generative space of one language onto the different generative space of another, inevitably losing features of the original in the process. Yet translation has historically expanded the effective audience of ideas rather than homogenizing them,

by making diverse ideas accessible across language barriers that would otherwise prevent their communication.

18.2 The Structural Pattern

These historical cases share a common pattern. In each, a normalization mechanism applied to the communication channel reduced the surface diversity of the observable output without reducing the diversity of the ideas that output expressed. In some cases, normalization actively facilitated the dissemination of diverse ideas by reducing the cognitive barriers to comprehension across heterogeneous audiences.

The same structural pattern applies to LLM-mediated normalization. The novel features of LLMs—their computational power, their breadth across domains, their accessibility—may amplify both the normalization effects and the facilitating effects. Which of these amplifications dominates is a contingent empirical question, not one that can be settled by theoretical analysis alone.

19 Reexamining Recent Empirical Claims

Recent discussions of linguistic and cognitive homogenization have been articulated most explicitly in the work of Sourati et al. [2026], who argue that large language models may gradually standardize human expression, perspective, and reasoning. Their analysis raises important concerns regarding the cultural consequences of generative AI systems and has stimulated a broader scholarly conversation about the dynamics of human–machine communication.

The information-theoretic and ecological framework developed in the preceding sections provides several grounds for reconsidering specific assumptions that underlie their claims.

19.1 Representation Versus Generation

A central premise of the homogenization argument is that similarities in linguistic output reflect convergence in the underlying cognitive processes that produced those outputs. As the formal analysis of Section 3 demonstrates, this premise is false as a general logical claim. A normalization layer can produce output similarity from diverse inputs; observing output similarity is therefore not evidence of input similarity.

For the premise to hold, it would need to be established that the normalization layer in question is approximately injective—that different cognitive processes tend to produce distinctly different model-assisted outputs. This claim requires positive evidence, not merely the assumption that output tracks input.

19.2 Instrumental Versus Epistemic Interaction

The homogenization argument often treats model outputs as epistemic influences that substantively reshape users’ conceptual frameworks. The distinction developed in Section 5 shows that this characterization applies only to epistemic modes of interaction. In instrumental modes, the model functions as a communication interface, and its influence on underlying reasoning is limited.

Existing empirical studies of homogenization have not, in general, adequately controlled for the difference between these modes. Studies that expose participants to opinionated model-generated content are studying the epistemic interaction mode, and their results cannot be generalized to instrumental use cases.

19.3 Short-Term Effects and Long-Term Dynamics

The empirical studies cited in support of homogenization typically employ short-term experimental designs: participants interact with model outputs during a single session, and subsequent shifts in expressed opinion or linguistic style are measured. Such designs are capable of detecting immediate alignment effects but provide no evidence about whether those effects persist once the interaction ends.

The relevant claim in the homogenization thesis is about long-term cultural dynamics: a gradual and durable shift in the generative processes of a population. Demonstrating such a shift requires longitudinal designs with appropriate washout periods and measures that can distinguish persistent changes in generative behavior from temporary surface alignment. The existing literature provides no such evidence.

19.4 Stylistic Similarity and Conceptual Diversity

Many studies supporting the homogenization thesis focus on stylistic or lexical features of text: the distribution of word frequencies, syntactic constructions, or rhetorical organizational patterns. These are precisely the features that normalization layers are designed to standardize. As noted above, observing stylistic convergence in model-assisted texts is consistent with—and indeed predicted by—a channel normalization account.

The relevant measure for the homogenization thesis is not stylistic similarity but *conceptual diversity*: the diversity of ideas, arguments, and interpretive frameworks expressed in the texts. Measuring conceptual diversity is substantially more difficult than measuring stylistic similarity, but it is the measure that the thesis actually requires.

20 Toward a Precise Research Agenda

The preceding critiques identify specific methodological gaps in the existing empirical literature. Filling these gaps requires a research agenda organized around a clearer conceptual framework. We propose three priority directions.

20.1 Longitudinal Studies of Generative Behavior

The central empirical question is whether sustained use of LLMs produces durable changes in the generative processes of human authors—changes that persist in the absence of model assistance. Addressing this question requires longitudinal study designs that track the same individuals over extended periods, with periodic measurement of unassisted writing under controlled conditions.

Such studies should measure not only surface stylistic features but also conceptual properties: the range of arguments made, the diversity of sources cited, the variety of framing strategies employed. If generative diversity contracts over time in unassisted writing, and if this contraction correlates with the intensity of model use, the homogenization thesis would receive genuine empirical support.

20.2 Transfer Studies in Reasoning

The reasoning convergence hypothesis requires evidence of durable transfer from model-mediated reasoning contexts to independent problem-solving behavior. Experimental designs could expose participants to extended interaction with model-generated reasoning traces for specific problem types and then measure problem-solving strategies on novel instances of the same types in the absence of model assistance. If transfer occurs and if the transferred strategies resemble model-generated strategies more than alternative human strategies for the same problems, this would constitute evidence for reasoning convergence.

20.3 Measurement of Exploratory Entropy

The concept of exploratory entropy introduced in Section 16 suggests a third research direction: direct measurement of how algorithmic mediation reshapes the distribution of sources that individuals encounter. Digital trace data from browsing behavior, citation networks, and search interaction logs could in principle be used to estimate $P(s)$ before and after the introduction of model-mediated search into a given domain.

Comparing the entropy of discovery distributions across populations with varying degrees of model-mediated versus direct-search engagement would provide direct evidence about the ecological effects of generative AI on knowledge exploration.

21 Cognitive Prostheses and Human Agency

Generative AI systems are most productively understood within the broader theoretical tradition of cognitive prostheses: technologies that externalize aspects of cognitive processing and thereby reconfigure the functional boundaries of the human cognitive system.

21.1 The Prosthetic Tradition

Writing itself is the paradigm case of a cognitive prosthesis. By externalizing memory into inscribed symbols, writing extended the temporal and spatial reach of human cognition in ways that fundamentally altered the structure of intellectual culture. Yet no one concludes from the universality of written language that human cognition has homogenized; rather, writing created new conditions for the proliferation of diverse ideas by enabling their stable transmission across time and space.

Printing extended this effect by distributing written content to larger audiences at lower cost. Calculators externalized arithmetic processing, freeing cognitive resources for higher-order mathematical reasoning. Search engines externalized the process of bibliographic exploration, enabling access to larger corpora than any individual could maintain in memory.

Each of these prosthetic technologies reshaped the landscape of cognitive activity without eliminating human agency in directing that activity. They introduced new affordances and new constraints, shifted the allocation of cognitive effort, and created new dependencies. But the direction of inquiry, the selection of questions, and the evaluation of conclusions remained irreducibly human.

21.2 LLMs as Prosthetic Systems

Large language models extend this lineage into new domains. They externalize substantial portions of the linguistic production process: drafting, editing, reformulation, summarization. They also externalize portions of the information retrieval and synthesis process: answering questions, explaining concepts, generating literature reviews.

The empirically relevant question is not whether this externalization influences human cognitive behavior—it surely does, as do all prosthetic technologies—but how it reshapes the space of available cognitive strategies. The history of prior prostheses suggests that the answer to this question is rarely simple and often counterintuitive. Calculators did not reduce mathematical thinking; they shifted it from arithmetic to conceptual domains. Search engines did not reduce reading;

they changed the patterns of reading. The effects of LLMs on human cognition are likely to exhibit comparable complexity.

21.3 Human Agency in the Prosthetic System

A central feature of prosthetic technologies is that they do not merely constrain the human who uses them but create new possibilities for action. The writer using a word processor has access to revision strategies that would be prohibitively costly with pen and paper. The researcher using a search engine can explore connections across bodies of literature that would be inaccessible within a physical library’s limits.

LLMs similarly create new possibilities: the ability to quickly prototype arguments, to translate idiomatic expression into standard forms, to generate alternative framings of a familiar problem. Whether individuals exploit these possibilities in ways that expand or contract the diversity of their intellectual output is a function of how they choose to use the technology—a function, in other words, of human agency operating within the new landscape the technology creates.

This observation has normative implications for the homogenization debate. If the cultural consequences of LLMs depend substantially on how they are used, the appropriate response to concerns about homogenization is not merely to analyze the statistical properties of model outputs but to design usage contexts, interface features, and pedagogical practices that encourage expansive rather than contractive modes of engagement.

22 Implications for Linguistic Diversity

The theoretical framework developed throughout this paper permits a more nuanced and precise account of LLMs’ effects on linguistic diversity than the homogenization thesis provides. Several implications warrant explicit statement.

22.1 Normalization and Diversity Are Compatible

The central formal result is that normalization in the communication channel is compatible with the full preservation of generative diversity at the source. Proposition 3.2 establishes this compatibility in the general case, and the historical cases of Section 18 provide empirical corroboration. There is no logical pathway from the observation of normalized outputs to the conclusion of cognitive homogenization.

22.2 Normalization May Expand Effective Expressive Space

The argument of Section 4 suggests that automated normalization may expand rather than contract the effective expressive space of individual authors by removing the communicative penalty that previously attached to non-standard expression. This possibility is underexplored in existing literature and merits empirical investigation.

22.3 The Ecological Risk Is Real but Differently Located

The framework identifies a genuine cultural risk associated with LLM-mediated communication, but locates it differently from the homogenization thesis. The risk lies not in the normalization of expressive form but in the potential contraction of exploratory entropy: the diversity of pathways through which individuals encounter ideas. Propositions 16.4 and 17.1 formalize the conditions under which this contraction occurs and demonstrate that it is a structural feature of summarization and ranking systems rather than a contingent effect.

22.4 Design Choices Determine Cultural Consequences

The magnitude of the exploratory entropy reduction introduced by AI systems depends substantially on their design: the specificity of their summaries, the diversity of sources they surface, the degree to which they support rather than replace direct engagement with original materials, and the stochasticity of their discovery mechanisms. These are tractable design choices, not structural inevitabilities.

23 Externalization of Linguistic Effort and the Redistribution of Cognitive Load

Throughout the history of intellectual practice, technological systems have repeatedly altered the distribution of cognitive effort required to produce and transmit ideas. Writing externalized memory, allowing information to persist outside the human mind and enabling the systematic accumulation of knowledge across generations. Printing externalized the reproduction of texts, enabling ideas to circulate across large populations at low marginal cost. Calculators externalized arithmetic operations, releasing attentional resources for higher-level reasoning about structure and relationship. Search engines externalized bibliographic exploration, enabling rapid navigation across bodies of literature that would otherwise require years to traverse.

Large language models extend this lineage into new domains. They externalize substantial portions of the linguistic production process that previously required

significant cognitive effort from human authors: drafting, grammatical correction, stylistic reformatting, translation between registers, and summarization. When such tasks are externalized, the total cognitive resources available to an author do not necessarily decline; they are redistributed across different stages and kinds of intellectual work.

23.1 Reallocation of Compositional Resources

When linguistic optimization becomes inexpensive, the marginal cost of expressing complex or speculative ideas decreases. Authors may become more willing to pursue unconventional arguments, extended analogies, and interdisciplinary syntheses when the linguistic burden of expressing those ideas is substantially reduced. Ideas that would previously have been abandoned partway through composition—because the effort required to formulate them clearly exceeded the available cognitive budget—may now reach expression more frequently.

The redistribution also alters the temporal structure of composition. Exploratory drafting practices become less costly when the mechanical work of rendering provisional expressions into polished form is partly automated. Authors can generate approximate articulations of ideas rapidly, relying on the normalization system to refine the result and encouraging iteration across a wider range of conceptual possibilities rather than deeper development of a narrower set.

This effect is an extension of the well-documented impact of word processing on writing practice. Studies of the transition from typewriters and handwritten manuscripts to word processors found that the reduction in revision costs encouraged longer, more complex, and more iteratively developed documents. The cognitive resources freed from mechanical drafting were partially redirected toward substantive revision of content. Automated normalization may produce a comparable effect at the level of stylistic and grammatical refinement.

23.2 The Separation of Generation and Representation

Externalization of linguistic effort introduces a more explicit structural separation between the generation of conceptual content and the production of its public representation. The generative stage of writing can proceed with greater tolerance for incomplete, irregular, or idiosyncratic expression when the representational stage is handled by an automated system. This division resembles the separation between drafting and editing that exists in collaborative writing environments—the difference being that the editor is now automated and available at negligible cost.

From the perspective of the homogenization debate, the externalization of linguistic effort introduces an important asymmetry. The processes most directly affected by language models are those associated with representation and formatting

rather than those associated with the generation of ideas. The reduction of stylistic variation in final outputs reflects a shift in how linguistic labor is distributed, not a contraction of conceptual possibility. The resources previously devoted to achieving linguistic conformity are released for other purposes; the question is where they are redirected, and the evidence from analogous historical transitions suggests the answer is: toward substantive intellectual work.

23.3 Historical Precedents

The pattern of cognitive resource reallocation through technological externalization is well-established in the history of intellectual tools. The introduction of algebraic notation in mathematics freed mathematicians from tracking the referents of verbal descriptions and enabled the development of abstract structural reasoning that was previously too cognitively costly to sustain. The introduction of statistical software freed researchers from manual calculation and enabled the widespread adoption of more complex models. In each case, the externalization of a cognitive operation did not reduce the quality or diversity of the intellectual work it supported; it changed the level at which that work was conducted.

The externalization of linguistic production through LLMs is likely to produce comparable effects: a shift in the level at which human intellectual effort is concentrated, from the surface properties of expression toward the structural and conceptual properties of argument. Whether this shift will on balance increase or decrease intellectual diversity is an empirical question. But the theoretical analysis suggests there is no strong reason to expect it to decrease diversity, and several reasons to expect the opposite.

24 Temporal Signatures in Intellectual Production

Intellectual identity is not defined solely by the spatial distribution of topics within conceptual space. It also emerges through temporal patterns in the production and development of ideas. Over extended periods, authors exhibit characteristic rhythms in the way they explore, refine, and recombine conceptual materials. These rhythms form temporal signatures that remain visible across bodies of work even when the linguistic surface of individual texts becomes standardized.

24.1 Temporal Structure of Conceptual Trajectories

Let $\{x_1, x_2, \dots, x_n\}$ denote the sequence of texts produced by an author over time, ordered by date of composition. The corresponding topic vectors $(v_1, v_2, \dots, v_n)$ in Δ^{m-1} represent a time-stamped trajectory through conceptual space. The structure

of this trajectory is not fully captured by the set of visited positions; the ordering and timing of those positions carry additional information about the dynamics of the author’s intellectual process.

The first-order difference sequence $\Delta v_k = v_{k+1} - v_k$ describes the local direction of movement through conceptual space at each step. Sequences of such directional changes encode the characteristic rhythms of intellectual exploration: the pace at which new topics are introduced, the tendency to consolidate or diverge after a period of synthesis, the characteristic duration of engagement with a particular framework before moving on.

These temporal structures have been observed in many historical cases. Scientific research programs often progress through recognizable phases: a period of exploratory experimentation during which the problem space is mapped, a consolidation phase during which theoretical frameworks are developed, and an extension phase during which the framework is applied to new domains. Philosophical traditions exhibit comparable cycles of critique, reformulation, and synthesis. Literary and artistic careers show characteristic developmental arcs that manifest in the temporal structure of thematic evolution.

24.2 Temporal Signatures as Author Fingerprints

Temporal patterns in conceptual trajectories constitute a form of intellectual fingerprint that is orthogonal to both stylistic and purely distributional characterizations of authorship. Two authors may have similar time-averaged topic distributions while exhibiting very different temporal dynamics: one moving rapidly between conceptual domains, the other developing a single framework in depth over an extended period.

These temporal structures are largely independent of the stylistic features of linguistic representation. A conceptual transition from one region of the intellectual landscape to another may be expressed through standardized prose without obscuring the underlying movement in conceptual space. The normalization of linguistic form leaves the temporal dynamics of intellectual development intact.

24.3 Temporal Records Under Reduced Friction

The availability of automated linguistic assistance may make temporal signatures more visible rather than less. When the mechanical effort of editing and formatting text decreases, the cost of producing intermediate drafts and exploratory notes also decreases. Authors may record more stages of their conceptual development, producing richer temporal records of their intellectual activity. The temporal trajectory—which was previously partially obscured by the selection pressure that

determined which drafts were worth the effort of production— may become more fully observable as that pressure decreases.

This suggests a counterintuitive consequence: automated normalization may increase the temporal resolution at which intellectual trajectories can be studied, by reducing the costs that previously determined how much of the trajectory was recorded. The cultural consequence would be an expansion of the observable record of intellectual diversity in time, not a contraction.

25 A Multilayer Model of Expressive Diversity

The preceding sections have introduced several distinct dimensions of diversity within AI-mediated communication systems. Each corresponds to a different layer in the generative and communicative process through which ideas are produced and transmitted. This section synthesizes these layers into a unified model and draws out the implications for evaluating the homogenization thesis.

25.1 The Five Layers

At the deepest level lies the diversity of generative conceptual processes, reflected in the source entropy $H(X)$. This layer corresponds to the variety of conceptual configurations that individuals are capable of producing—the combinatorial expansion described in Section 14 ensures that this space remains effectively unbounded.

Above this lies the diversity of conceptual trajectories and toolchains. Authors move through the conceptual manifold in distinctive patterns, forming recognizable constellations of topics and recurring sequences of analytical procedures. The topology of these trajectories—including the occurrence of transverse movements between conceptual strata—encodes stable intellectual signatures that persist across bodies of work (Sections 8 and 9).

A third layer consists of behavioral signals associated with the embodied production of communication. Voice patterns, handwriting trajectories, typing dynamics, and other input modalities introduce additional degrees of behavioral variation that originate in processes that are invisible to and unaffected by linguistic normalization systems (Section 10).

A fourth layer corresponds to temporal signatures: the characteristic rhythms at which conceptual positions evolve over time, the dynamics of exploration and consolidation, and the developmental arcs through which intellectual practices mature (Section 24).

The fifth and outermost layer consists of the observable linguistic representation transmitted through the communication channel. The entropy $H(Y)$ of this representation may be reduced by normalization processes, and the exploratory entropy

H_E of the discovery pathways through which ideas are encountered may be reduced by algorithmic mediation (Sections 16 and 17).

25.2 Layer Independence and the Invalidity of Homogenization Inferences

These layers are only weakly coupled. Compression at the level of linguistic representation does not propagate automatically to the deeper layers of conceptual generation, behavioral expression, or temporal dynamics. A communication system may exhibit substantial uniformity at its surface output while maintaining considerable diversity in the processes that generate those outputs.

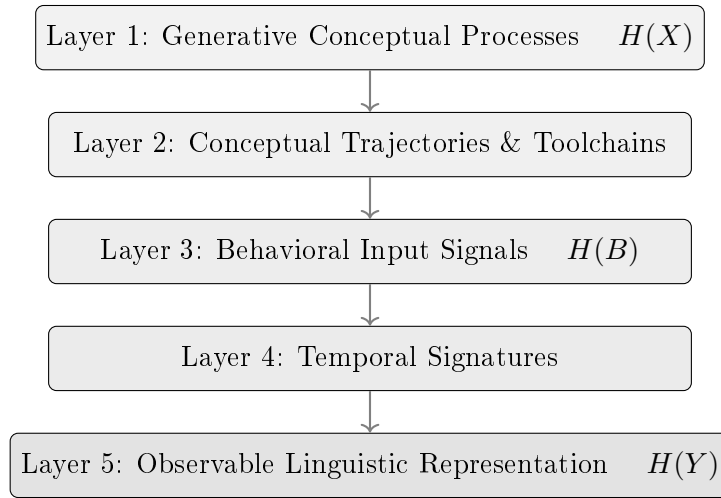


Figure 2: Five-layer model of expressive diversity in AI-mediated communication. LLMs as normalization systems act primarily at Layer 5; the deeper layers are governed by independent mechanisms and are not directly constrained by surface representation.

The homogenization thesis implicitly assumes that reductions in observable linguistic entropy $H(Y)$ correspond to reductions in the generative entropy $H(X)$ of the underlying cognitive system. The multilayer model demonstrates that this inference is unwarranted. Each layer possesses its own dynamics, and changes at one layer need not imply corresponding changes at another.

Large language models primarily affect the outermost representational layer. They rewrite and normalize textual expressions without directly constraining the conceptual configurations that generate those expressions, the behavioral patterns through which they are produced, the temporal dynamics through which they evolve, or—as argued separately—the exploratory pathways through which source material is encountered. At the same time, their influence on discovery pathways may reshape the distribution of intellectual inputs that authors encounter, introducing a distinct and more tractable form of cultural concern.

25.3 Implications for Research and Policy

The multilayer model reshapes the research agenda for evaluating the cultural consequences of generative AI. Studies that measure only the properties of the observable output layer—stylometric analyses of LLM-assisted texts, experiments on short-term opinion alignment, distributional comparisons of model-generated prose—are measuring at the wrong level. They capture the behavior of the normalization transformation, not the behavior of the generative cognitive system.

Research designs sensitive to the deeper layers would require longitudinal tracking of conceptual trajectories, comparison of toolchain distributions across corpora with varying levels of AI assistance, measurement of exploratory entropy in digital knowledge environments, and experimental investigation of how externalized linguistic labor reshapes the allocation of cognitive effort.

Policy interventions aimed at preserving intellectual diversity should similarly target the layers where the genuine risks are concentrated: the exploratory infrastructure through which ideas are discovered, the diversity of conceptual frameworks surfaced by summarization systems, and the design of interfaces that determine whether AI tools expand or constrain the range of intellectual possibilities available to users.

26 Conclusion: Diversity, Compression, and Cultural Dynamics

The homogenization thesis rests on the assumption that uniformity in linguistic outputs implies convergence in human cognition. This assumption is formally invalid: a normalization layer in the communication channel can produce the former without the latter. The data processing inequality—that post-processing cannot increase mutual information—establishes this result as a general theoretical claim, not merely a possibility in specific cases.

Large language models, when used instrumentally as editing, translation, or drafting tools, operate as normalization layers. They map diverse generative inputs onto a more standardized output distribution. The entropy of that output is necessarily lower than the entropy of the source distribution; but this reduction is a property of the transformation, not evidence that the source distribution has changed. Observing stylistic convergence in LLM-assisted texts no more establishes cognitive convergence than observing orthographic uniformity in printed texts establishes uniformity in the ideas those texts express.

The deeper analysis developed in this paper identifies several additional dimensions along which expressive and intellectual diversity persists beneath surface linguistic uniformity. Conceptual trajectories through the topic simplex Δ^{m-1} re-

main distinctive across authors even when their linguistic surfaces converge; topic constellations and conceptual toolchains form stable intellectual signatures that grow more precisely identifiable with larger bodies of work; behavioral input signals from embodied production processes introduce high-dimensional variation that is invisible to textual normalization; temporal dynamics in intellectual production encode characteristic rhythms of exploration and consolidation; and the combinatorial expansion of conceptual space ensures that the number of intellectually novel configurations remains effectively unbounded.

The genuine cultural risk identified by this analysis lies elsewhere: in the potential contraction of exploratory entropy as knowledge discovery becomes increasingly centralized and algorithmically mediated. The diversity of conceptual pathways through which individuals encounter ideas—a significant source of intellectual innovation and cultural resilience—may contract as a consequence of the summarization and ranking dynamics that characterize current AI systems. This is a distinct phenomenon from linguistic normalization, governed by distinct mechanisms, and requiring distinct empirical and normative responses.

This is a tractable policy and design problem. Interfaces that preserve exploratory diversity, that surface conflicting interpretations alongside canonical ones, that encourage direct engagement with original sources, and that expose users to unexpected conceptual neighborhoods can maintain higher exploratory entropy than those that optimize exclusively for efficiency and convenience. The cultural consequences of generative AI depend on these choices in ways that are amenable to deliberate intervention.

The multilayer model synthesized in Section 25 provides the appropriate framework for this intervention. Intellectual diversity is distributed across five layers of the communication system—generative processes, conceptual trajectories, behavioral signals, temporal dynamics, and linguistic representation—and the appropriate measures of its preservation differ at each level. Understanding the long-term implications of generative language technologies requires sustained attention to all five layers: the diversity of generative sources, the structure of conceptual exploration, the behavioral richness of embodied expression, the temporal shape of intellectual development, and the ecology of knowledge discovery through which ideas circulate. Conflating these dimensions—as the homogenization thesis characteristically does—produces misdirected concern and obscures the places where cultural vigilance is most needed.

Appendices

A Formal Derivation of Entropy Inequalities Under Deterministic Normalization

This appendix collects and extends the entropy-theoretic results discussed in the main text.

A.1 Setup

Let $(\mathcal{X}, p)$ be a discrete probability space representing the generative source, where $\mathcal{X}$ is the set of possible linguistic expressions and $p(x) = \mathbb{P}(X = x)$ is the source distribution. Let $\mathcal{Y}$ be the output space and $T : \mathcal{X} \rightarrow \mathcal{Y}$ a deterministic function representing the normalization transformation. The observable output is the random variable $Y = T(X)$, with distribution

$$p(y) = \mathbb{P}(Y = y) = \mathbb{P}(T(X) = y) = \sum_{\{x: T(x)=y\}} p(x).$$

A.2 Entropy of the Source

$$H(X) = - \sum_{x \in \mathcal{X}} p(x) \log p(x).$$

A.3 Entropy of the Normalized Output

$$H(Y) = - \sum_{y \in \mathcal{Y}} p(y) \log p(y) = - \sum_{y \in \mathcal{Y}} \left(\sum_{\{x: T(x)=y\}} p(x) \right) \log \left(\sum_{\{x: T(x)=y\}} p(x) \right).$$

A.4 The Core Inequality

Theorem A.1 (Entropy Reduction Under Deterministic Normalization). *Let X be a discrete random variable and $Y = T(X)$ a deterministic function of X . Then*

$$H(Y) \leq H(X),$$

with equality if and only if T is injective on the support of $p(x)$.

Proof. Since Y is a deterministic function of X , the conditional entropy satisfies $H(Y | X) = 0$. By the chain rule:

$$H(X, Y) = H(X) + H(Y | X) = H(X).$$

Also by the chain rule:

$$H(X, Y) = H(Y) + H(X | Y).$$

Therefore $H(X) = H(Y) + H(X | Y)$, so $H(Y) = H(X) - H(X | Y) \leq H(X)$ since $H(X | Y) \geq 0$.

For equality: $H(X | Y) = 0$ if and only if X is a function of Y , which holds if and only if T is injective on the support of $p(x)$. $\square$

A.5 Quantifying Information Loss

The information lost in normalization is measured by the equivocation $H(X | Y)$:

$$H(X | Y) = H(X) - H(Y) = \sum_{y \in \mathcal{Y}} p(y) H(X | Y = y),$$

where

$$H(X | Y = y) = - \sum_{\{x: T(x)=y\}} \frac{p(x)}{p(y)} \log \frac{p(x)}{p(y)}.$$

This conditional entropy measures the residual diversity within each equivalence class $T^{-1}(y)$. When normalization collapses many diverse expressions into a single output, $H(X | Y = y)$ is large for that y , reflecting the substantial generative diversity that is invisible in the normalized representation.

A.6 The Data Processing Inequality

The main theorem is a consequence of the more general data processing inequality. If Z is any random variable such that $X \rightarrow Y \rightarrow Z$ is a Markov chain (i.e., Z depends on X only through Y), then

$$I(X; Z) \leq I(X; Y),$$

where $I(X; Y) = H(X) - H(X | Y)$ denotes mutual information. For deterministic $Y = T(X)$, $I(X; Y) = H(Y)$, and the inequality reduces to our main result.

B Exploratory Entropy: Formal Treatment

B.1 Setup

Let $\mathcal{S} = \{s_1, \dots, s_n\}$ be a finite set of informational sources and P a probability distribution over $\mathcal{S}$ representing the discovery probability. The exploratory entropy

is

$$H_E(P) = - \sum_{i=1}^n P(s_i) \log P(s_i).$$

B.2 Maximum Entropy and the Uniform Distribution

By the maximum-entropy principle, $H_E(P)$ is uniquely maximized by the uniform distribution $P_0(s_i) = 1/n$, at which point $H_E(P_0) = \log n$.

B.3 Entropy Under Ranking

Let $R : \mathcal{S} \rightarrow \{1, \dots, n\}$ be a bijective ranking function and $f : \{1, \dots, n\} \rightarrow (0, \infty)$ a strictly decreasing exposure function. Define the rank-weighted distribution:

$$P'(s_i) = \frac{f(R(s_i))}{\sum_{j=1}^n f(R(s_j))}.$$

Since f is strictly decreasing, P' is nonuniform, and therefore $H_E(P') < \log n = H_E(P_0)$.

Lemma B.1 (Rank Concentration). *For any strictly decreasing f , the rank-mediated distribution P' satisfies*

$$H_E(P') < H_E(P_0).$$

Moreover, as f becomes more steeply decreasing (as $f(1)/f(n)$ increases), $H_E(P')$ decreases monotonically toward zero.

Proof. The first claim follows from the strict concavity of the entropy function on the probability simplex: the unique maximizer of H_E is P_0 , and $P' \neq P_0$ for any nonconstant f . The second claim follows from the observation that as $f(1)/f(n) \rightarrow \infty$, all probability mass concentrates on the top-ranked source, and $H_E(P') \rightarrow 0$. $\square$

B.4 Entropy Under Summarization

Let $G : \mathcal{P}(\mathcal{S}) \rightarrow \mathcal{S}_*$ be a generative summarization operator that maps subsets of $\mathcal{S}$ to a summary space $\mathcal{S}_*$. If $\mathcal{S}_* = \{*\}$ (a single canonical summary), the discovery distribution degenerates to a point mass with $H_E = 0$.

In the realistic mixed case where users access the summary with probability α and access original sources (distributed according to P') with probability $1 - \alpha$, the effective discovery distribution is the mixture

$$P_\alpha = \alpha \cdot \delta_* + (1 - \alpha) \cdot P',$$

where δ_* is the point mass on the summary. The entropy of this mixture satisfies

$$H_E(P_\alpha) \leq (1 - \alpha)H_E(P') + h(\alpha),$$

where $h(\alpha) = -\alpha \log \alpha - (1 - \alpha) \log(1 - \alpha)$ is the binary entropy function. This bound shows that increasing α (greater reliance on AI summaries) monotonically reduces the upper bound on exploratory entropy.

C A Unified Proposition on Normalization and Exploration

Proposition C.1 (Multilayer Independence of Expressive Diversity). *The following hold simultaneously:*

- (i) *Normalization layers may reduce observable linguistic entropy $H(Y)$ while leaving generative cognitive entropy $H(X)$ unchanged (Theorem A.1).*
- (ii) *Algorithmic discovery systems may reduce exploratory entropy H_E even when both $H(X)$ and $H(Y)$ remain unchanged (Lemma B.1).*
- (iii) *Linguistic normalization does not constrain conceptual trajectory structure in Δ^{m-1} : distinct conceptual toolchains can produce similar normalized outputs, but the converse inference from output similarity to toolchain similarity is not warranted.*
- (iv) *The space of intellectually innovative configurations grows exponentially with $|\mathcal{T}|$, independently of $H(Y)$.*
- (v) *Consequently, $H(X)$, $H(Y)$, H_E , and conceptual trajectory structure are mutually logically independent variables in AI-mediated communication systems.*

Proof. Claims (i) and (ii) follow from Theorem A.1 and Lemma B.1 respectively.

For claim (iii): the mapping $F : \mathcal{C} \rightarrow \mathcal{L}$ is many-to-one, and the composition $T \circ F$ compounds this non-injectivity. Therefore $T(F(C_1)) = T(F(C_2))$ is consistent with $C_1 \neq C_2$ for arbitrary pairs of conceptual configurations.

For claim (iv): the number of subsets of $\mathcal{T}$ is $2^{|\mathcal{T}|}$, which grows exponentially in $|\mathcal{T}|$ regardless of the value of $H(Y)$.

For claim (v), independence follows from the construction of explicit examples:

High $H(X)$, low $H(Y)$, low H_E , homogeneous trajectories: Diverse authors use the same AI editor and same AI-mediated discovery channel, drawing from the same canonical sources.

High $H(X)$, low $H(Y)$, high H_E , diverse trajectories: Diverse authors use the same AI editor but explore independently through physically distributed archives.

High $H(X)$, high $H(Y)$, low H_E : Authors with diverse expression and minimal AI editing all rely on the same AI-curated news feed, producing diverse prose about a small set of concepts.

Low $H(X)$, low $H(Y)$, high H_E : Authors with narrow stylistic range and heavy AI editing nevertheless encounter ideas through a diverse and decentralized discovery network. $\square$

Remark C.2. This proposition clarifies the structure of the homogenization debate. The three quantities $H(X)$, $H(Y)$, and H_E capture three distinct dimensions of diversity in an AI-mediated knowledge system: the diversity of generative cognitive processes, the diversity of observable linguistic outputs, and the diversity of epistemic discovery pathways. Each is a meaningful and distinct object of study. Conflating them—as the homogenization thesis tends to do—produces a misdirected research program that measures the wrong variables and draws unwarranted conclusions.

References

- C. E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27:379–423, 1948.
- T. M. Cover and J. A. Thomas. *Elements of Information Theory*. Wiley-Interscience, 2nd edition, 2006.
- R. Boyd and P. J. Richerson. *The Origin and Evolution of Cultures*. Oxford University Press, 2005.
- J. Henrich. *The Secret of Our Success: How Culture Is Driving Human Evolution, Domesticating Our Species, and Making Us Smarter*. Princeton University Press, 2015.
- D. Lazer, A. Pentland, L. Adamic, S. Aral, A. Barabasi, D. Brewer, N. Christakis, N. Contractor, J. Fowler, M. Gutmann, T. Jebara, G. King, M. Macy, D. Roy, and M. Van Alstyne. Computational social science. *Science*, 323(5915):721–723, 2009.
- A. Sourati, J. Kim, T. Westberg, and M. T. Corey. The homogenizing effect of large language models on human expression and thought. *Trends in Cognitive Sciences*, 30(2):112–128, 2026.
- E. M. Bender, T. Gebru, A. McMillan-Major, and S. Shmitchell. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, pages 610–623, 2021.
- M. Raghavan and K. Hanna. Mechanisms of algorithmic influence on human decisions. *Annual Review of Economics*, 15:327–362, 2023.
- E. L. Eisenstein. *The Printing Press as an Agent of Change*. Cambridge University Press, 1979.
- A. Clark. *Being There: Putting Brain, Body, and World Together Again*. MIT Press, 1997.
- E. Hutchins. *Cognition in the Wild*. MIT Press, 1995.
- D. R. Olson. *The World on Paper: The Conceptual and Cognitive Implications of Writing and Reading*. Cambridge University Press, 1994.
- D. Sperber. *Explaining Culture: A Naturalistic Approach*. Blackwell, 1996.
- C. R. Sunstein. *#Republic: Divided Democracy in the Age of Social Media*. Princeton University Press, 2017.